How Residual-Stream Updates Travel Through Depth

A Pre-Registered Study of Cross-Layer Dynamics in Transformer Factual Recall

Aljaž Frančič
Independent Researcher
aljaz.francic@gmail.com

Draft of July 14, 2026

Abstract

Transformer components communicate through an additive residual stream, so the update a layer writes, $\Delta_\ell = v_{\ell+1} - v_\ell$, is the natural unit of cross-layer computation. We pre-registered two hypotheses about these updates on factual-recall inputs — (H1) they are predictable across depth; (H2) the prediction kernel is approximately depth-stationary — and ran the frozen protocol on Pythia-410m-deduped, replicating on five further models from three families. **(1)** H1 holds far above pre-registered nulls but decomposes exactly: an orthogonal split of the held-out explained variance attributes most of it to the running-sum tautology on the four deepest updates ($0.65 = 0.52 + 0.13$) and roughly half on the whole stack (raw $0.52 = 0.27 + 0.25$), and on the primary model what survives is substantially architectural — an untrained network exceeds it (though on GPT-2-small it is about half learned). **(2)** H2 fails decisively: kernels fit on one depth band predict another band's updates worse than that band's mean, so the right abstraction is a first-order ($k^* \approx 1$), fully depth-indexed state-space model (SSM), validating the unconstrained parameterisation of cross-layer transcoders. **(3)** Fitted over frozen sparse sources, the SSM reconstructs Δ_ℓ above a matched per-layer dictionary (+0.049 R^2 rank-matched, +0.008 parameter-matched), with a depth-graded, roughly half-learned gain; its gauge-invariant state path attenuates by half per layer of depth gap (memory $\approx$1.2 layers, consistent with $k^* \approx 1$) and the state decodes the recalled answer at mid-depth better than dimension-matched PCA of its own codes. **(4)** The pre-registered causal test returns a controlled null: no single feature's ablation clears threshold, and the group effect that does is reproduced by a norm-matched *random direction* — recall is a distributed, magnitude-driven code. As a by-product, on the primary model TopK dictionaries make single-feature directions seed-unidentifiable where matched L1 dictionaries on the same activations are perfectly stable. All gates were frozen before data collection, every headline reconstruction gap carries a paired bootstrap confidence interval, and every figure regenerates from tracked tables.

1 Introduction

Mechanistic interpretability decomposes transformer activations mostly layer by layer. Sparse autoencoders (SAEs) recover human-legible features from a single layer's residual stream (Cunningham et al., 2024; Bricken et al., 2023; Templeton et al., 2024), but a per-layer dictionary answers only a per-layer question — *what features live here?* — and not the question that matters for circuit analysis and causal editing: *how does what a layer writes relate to what earlier layers wrote?* The field has begun to attack this directly. Crosscoders fit one dictionary across all layers (Lindsey et al., 2024); multi-layer SAEs (MLSAEs) find that the same latent fires at different depths on different prompts (Lawson et al., 2025); and cross-layer transcoders (CLTs), the substrate of attribution-graph circuit tracing, give every feature a *free* decoder at every downstream layer — $O(L^2)$ independent decoder matrices with no prior connecting them (Ameisen et al., 2025).

This paper asks whether that free-parameter budget is load-bearing or redundant, and it asks the question the way one asks it in signal processing: by modelling the *increments*. Because the residual stream is additive, $v_{\ell+1} = v_\ell + \Delta_\ell$, regressing states on states trivially recovers the identity; the non-trivial object is the update $\Delta_\ell = v_{\ell+1} - v_\ell$ — what the layer *adds* — exactly as differencing removes the trend in classical time-series analysis. We pre-registered two nested hypotheses about these updates on factual-recall inputs:

H1 (depth-local predictability). Δ_ℓ is predictable from preceding updates $\Delta_{\ell-1}, \Delta_{\ell-2}, \ldots$ by a low-rank linear map, above scrambled-activation and context-perturbation nulls — residual updates form an autoregressive process along depth.

H2 (approximate stationarity). The prediction kernel is approximately shared across depth, so that a single block-Toeplitz (convolutive) kernel $\Delta_\ell \approx \sum_k W_k \, s_{\ell-k}$ over sparse per-layer sources is a useful generative model. Under H2, CLT's $O(L^2)$ decoder matrices would collapse to $O(K)$ shared kernel taps.

H2's form was motivated by an analogy to convolutive blind source separation in high-density surface electromyography, where sources mix into recordings through characteristic propagation kernels (Holobar and Zazula, 2007); the analogy generated the hypothesis and nothing else — the formal models and tests are native to interpretability and classical statistics. Crucially, the study was *pre-registered*: a frozen protocol fixed the diagnostics, the go/no-go gates, the null controls, the decomposition architecture as a function of the gate outcomes, and the causal-evaluation thresholds, before any data were collected. Where the data contradicted the more exciting branch, we took the pre-registered exit. Both hypotheses could have failed cheaply; what actually happened is more interesting than either clean outcome. Scope, stated up front: all kernels are fit at the *query position* of factual-recall prompts drawn from three CounterFact relation classes. The binary characterisation replicates on generic text, but kernel transfer fails across both token positions (eleven of twelve cross-position transfers negative) and relation classes (transfer ≈ 0 or negative; Section 4.3) — the kernels are hyper-local to their circuit, itself a finding — so every claim below is a claim about that regime.

What we found. H1 survives every control we could construct — including two nulls specified in advance and a post-hoc pressure battery designed to break it — but decomposes on inspection: an exact (Pythagorean) split of the headline autoregressive R^2 shows that on factual prompts the running-sum tautology takes the majority share on the pre-registered headline fit ($0.65 = 0.52 + 0.13$, scored on the four deepest updates) and roughly half on the whole-stack fit ($0.52 = 0.27 + 0.25$), and an untrained-network control shows the tautology-free remainder is, on the primary model, *substantially architectural* — a randomly initialised transformer exceeds it. Scored at a common lag order the same control makes GPT-2-small about half learned, so this is a Pythia result, not a law (Section 3). H2 fails decisively and non-negotiably: an AR kernel fit on one third of the depth axis predicts another third's updates *worse than that band's own mean*, at every rank, with no clean segment boundary (Section 4). The pre-registered consequence is that the generative model must be a first-order, fully depth-indexed state-space model (SSM) — and that CLT's refusal to share parameters across depth is *validated*, not wasteful.

Fitting that SSM over frozen per-layer SAE sources yields the paper's central positive result (Section 6): cross-layer state reconstructs residual updates above a matched per-layer dictionary at every sparsity ($+0.049$ R^2 $[+0.046, +0.052]$ at matched rank, $+0.008$ $[+0.005, +0.010]$ at matched total parameters), the gain grows monotonically with depth within Pythia (6 layers $+0.008 \rightarrow$ 24 layers $+0.050$; the sign, though not the trend, replicates at one point each in GPT-2 and Qwen2), and about half of it is present at random initialisation — so the method captures a genuinely learned signal roughly equal in size to the architectural one. The fitted model's gauge-invariant state path attenuates by half per layer of depth gap (implied memory ≈ 1.2

layers), consistent with — and supplying a candidate mechanism for — the observed first-order structure, and the state is functionally informative: at matched dimension it decodes the recalled answer at mid-depth better than PCA summaries of the codes it is built from.

To be direct about the framing: this is a *characterisation* paper, not a method paper. We do not advance the SSM as a better decomposition tool on the strength of the +0.008 parameter-matched margin — read alone, that number says most of the rank-matched gap is parameter count. The case that the cross-layer state captures real structure rests on what decoder width cannot buy at any budget: the growth of the gain with depth, its roughly-half-learned split against untrained controls, the gauge-invariant state-path decay, and the state's functional informativeness for the recalled answer. Those structural results, not the headline gap, carry the paper's positive claims — and the depth series carries its own confounds (N, width, dictionary type; Section 6.2), so we lean on it least.

The pre-registered causal test — find a feature no baseline finds, ablate it, predict the behavioural consequence — returns a null we consider one of the paper's main contributions (Section 6.6). No single feature clears the pre-registered effect threshold; ablating each example's top-100 co-firing features does clear it, but a *norm-matched random direction* reproduces that group effect, so it is a magnitude effect, not evidence of a privileged code. The null is dictionary-independent: a cleaner L1 dictionary concentrates single-feature effects nearly five-fold and still finds no bottleneck. Factual recall in this model is a distributed, mass-driven code.

Contributions.

1. A pre-registered, gate-structured characterisation of residual-update dynamics: H1 holds (with an exact accounting of how much is tautology, and evidence the remainder is substantially architectural on the primary model but roughly half learned on GPT-2-small; accounting and untrained control run on two of the six models); H2 fails (pervasive non-stationarity, $k^* \approx 1$, query-position-specific and circuit-specific kernels). The sign pattern replicates across six models, three families, and — on the primary model — on generic Pile text, where the effective order is k^* median 2 rather than 1 (Sections 3 and 4).

2. A first-order, fully depth-indexed SSM decomposition of Δ_ℓ over sparse per-layer sources, with a reconstruction gain over matched per-layer baselines that is CI-significant, depth-graded, and approximately half learned / half architectural (Sections 5 and 6).

3. A three-axis comparison against the unconstrained $O(L^2)$ residual-write CLT parameterisation: at matched per-decoder rank the CLT keeps a small significant edge; at matched *total* parameters the SSM Pareto-dominates up to 52M parameters; below ∼1k training examples the SSM generalises better (Section 6.4).

4. A clean, fully controlled causal null: recall is distributed and magnitude-driven, with no privileged sparse subset — established with matched-count, matched-mass, and norm-matched random-direction controls, and robust to the dictionary (Section 6.6).

5. Two structural positives that survive every robustness check: the fast-attenuating propagation path — the *gauge-invariant* state-path kernel halves per layer of depth gap, stable across seeds and dictionaries (we also show that the per-layer spectral radius, which an earlier version of this work reported as gauge-invariant, is not) — and the functional informativeness of the SSM state for answer identity at mid-depth (Sections 6.7 and 6.8).

6. A practical warning for the field: on the primary model, TopK dictionaries make single-feature decoder directions seed-unidentifiable (match@0.7 = 0.08 trained on the same activations, 0.45 for the public suite) where a matched L1 dictionary trained on those activations is perfectly stable (1.00) — an identifiability failure that lives in the sparsity mechanism, not the model or the decomposition (Section 6.5).

Everything reproduces from tracked run tables via one script; the protocol documents were frozen before data collection and every deviation is logged (Sections A and E).

2 Related work

The residual stream and superposition. Elhage et al. (2021) formalised the framing every cross-layer method inherits: attention and MLP blocks read from and write to a shared additive channel, so features persist through depth because the channel is a running sum. Elhage et al. (2022) explains why single neurons are not the unit of analysis: models encode more features than dimensions, in overlapping directions. Together these motivate both the object of this study, the per-layer write Δ_ℓ, and its principal statistical hazard — consecutive updates mechanically share a $\mp v_\ell$ term (Section 3.1), which we spend considerable effort separating from genuine structure.

Per-layer sparse dictionaries. SAEs trained on one layer's activations recover monosemantic features (Cunningham et al., 2024; Bricken et al., 2023) and scale to frontier models (Templeton et al., 2024). TopK SAEs replace the L1 penalty with a hard sparsity constraint (Makhzani and Frey, 2013; Gao et al., 2025). We use public per-layer suites as frozen sources — the MLSAE release for Pythia (Lawson et al., 2025), the reformatted GPT-2-small residual suite (Bloom, 2024), an open Qwen2 suite (Chanin, 2024), and EleutherAI's Pythia-70m-deduped suite (EleutherAI, 2024) — and additionally train matched L1 and TopK dictionaries on our own cached activations, which turns out to matter: the sparsity mechanism, not the model, controls whether single-feature decoder directions are identifiable across seeds (Section 6.5).

Cross-layer dictionaries. Crosscoders fit one dictionary with per-layer decoders across all layers at once (Lindsey et al., 2024). MLSAEs share the dictionary and let activations vary per layer; their headline observation — a latent is often active at exactly one layer per token, but at *different* layers across tokens — already strains the picture of a feature with a fixed home depth (Lawson et al., 2025). RouteSAE adds a router that picks which layer a shared SAE decodes per token (Shi et al., 2025): a *where-to-read* mechanism, complementary to the *how-updates-propagate* question here. Jacobian SAEs sparsify the causal sensitivities across a single MLP block (Farnik et al., 2025) — causal rather than correlational cross-layer structure, but with no assumption (or test) of depth-shared form.

Cross-layer transcoders. CLTs (Ameisen et al., 2025) are the closest prior method and this paper's primary conceptual baseline. A CLT feature reads at one source layer ℓ' and writes into every layer $\ell \geq \ell'$ through a separate, fully free decoder per (source, target) pair — $O(L^2)$ decoder matrices with no parameter sharing along depth. Our H2 is precisely the hypothesis that those matrices compress to a kernel in the depth difference $\ell - \ell'$; H2's failure (Section 4) is therefore an independent *validation* of CLT's unconstrained parameterisation, and our SSM-vs-CLT comparison (Section 6.4) asks the surviving quantitative question: how much reconstruction do the free per-pair decoders buy over a first-order depth-indexed state at matched sources?

Depth-wise lenses. The Tuned Lens learns per-layer affine probes that decode intermediate residual states into the final-output distribution (Belrose et al., 2023). We use it twice: its implied per-layer update $\hat{r}_{\ell+1} - \hat{r}_\ell$ is the calibrated baseline our cheap H1 diagnostic must beat before the expensive tests run (it does, Section 3.2), and it reappears as an update-prediction baseline in the evaluation, where — like every method not trained on Δ_ℓ — it scores below zero (Section 6.3).

Factual recall. We test on subject–relation–object recall because it is the best-instrumented circuit family in the literature: causal localisation to mid-depth MLPs (Meng et al., 2022; Dai et al., 2022), a staged enrich-then-extract account (Geva et al., 2023), and a standing caveat that causal-tracing locations and editing locations dissociate (Hase et al., 2023). Our data derive from CounterFact (Meng et al., 2022), with Pararel (Elazar et al., 2021) recorded in the protocol as the scaling path we ultimately did not need (Section A). Our distributed-code null (Section 6.6) is consistent with this literature's drift away from single-component stories of recall.

Signal-processing motivation. Convolution Kernel Compensation decomposes high-density surface EMG under a convolutive mixing model (Holobar and Zazula, 2007). The analogy — sources with characteristic propagation patterns along a depth-like axis — supplied H2's form (a shared kernel in the depth difference). It supplied nothing else: every model fit here is ridge regression, LASSO, alternating minimisation, or gradient descent on a linear SSM, and H2 was treated as a hypothesis to kill, which the data did.

Infrastructure. Activations are collected with TransformerLens (Nanda and Bloom, 2022) from Pythia (Biderman et al., 2023), GPT-2 (Radford et al., 2019), and Qwen2 (Yang et al., 2024) checkpoints; generic text is drawn from the Pile (Gao et al., 2020). Optimisation uses FISTA (Beck and Teboulle, 2009) for the sparse-code updates; the Granger-style tests are calibrated against a Freedman–Lane permutation scheme (Freedman and Lane, 1983) in Section C.

3 Phase 1 — is the residual update predictable across depth?

3.1 Study design

Pre-registration. The full protocol — four phases, per-step procedures, three go/no-go gates, null controls, decision matrices, and the causal-evaluation thresholds — was written and frozen before any data were collected. Progress and deviations were recorded exclusively in a separate tracker so that no pre-registered bar could drift. Every deviation from the plan (all quantitative-threshold adjustments explicitly sanctioned by the plan's own tuning rules, plus one descoping decision) is enumerated in Section A. The three gates: Gate 1 (is H1 supported above null?), Gate 2 (does H2 hold? — this picks the Phase-3 architecture), Gate 3 (does the decomposition reconstruct competitively?). A failed gate routes to writing up the negative result, not to silent continuation.

Primary model and notation. The primary model is Pythia-410m-deduped (Biderman et al., 2023): $L = 24$ blocks, residual dimension $D = 1024$, parallel attention/MLP blocks, with a public Tuned Lens and a public per-layer residual SAE suite. We write $v_\ell \in \mathbb{R}^D$ for the residual state *entering* block ℓ — equivalently, the running sum after the first ℓ blocks, with blocks indexed $0, \ldots, L-1$ (v_0 = embedding output; v_L = final pre-unembedding state, so 25 states per example) — and

$$\Delta_\ell \;=\; v_{\ell+1} - v_\ell, \qquad \ell \in \{0, \ldots, L-1\},$$

for the update block ℓ writes. Cross-model replications use Pythia-160m-deduped (L=12), GPT-2 small/medium/large (Radford et al., 2019) (L=12/24/36), and Qwen2-0.5B (Yang et al., 2024) (L=24) — with the primary, six models, three families, depths 12–36. Pythia-70m-deduped (L=6) is a seventh, debug-only lane: it serves as the optimiser-development model and the six-layer point of the depth-scaling series (Section 6.2), but its differently constructed dataset keeps it outside the H1/H2 replication claim (Section D).

Dataset. From CounterFact (Meng et al., 2022)[1] we keep triples the model actually knows: top-1 prediction equals the object and top-1 probability exceeds 0.1. Pythia-410m-deduped confidently knows only ∼2,000 of the 21,601 CounterFact triples, so the plan's tuning rules were exercised in the sanctioned direction (probability threshold 0.3 → 0.1; per-class floor 500 → 300), yielding **1,555 prompts in three relation classes** (organization 711, geography 511, language 333; occupation and religion fell below the floor and were dropped). Scoping to these three classes is Limitation 3 (Section 7). For every prompt we collect the residual trajectory $v_0, \ldots, v_{24}$ *at the final prompt token* — the query position, immediately before the model emits the object — giving a $[1555, 25, 1024]$ store. A fixed split (1,244 train / 311 test, seed 0) is shared by all Phase-1–3 analyses. Sanity check: applying the model's final LayerNorm and unembedding to the collected v_L recovers the model's top-1 answer (5/5 spot checks).

The $\pm v_\ell$ tautology, stated up front. Consecutive updates mechanically overlap: $\Delta_\ell = v_{\ell+1} - v_\ell$ and $\Delta_{\ell-1} = v_\ell - v_{\ell-1}$ share the term $\mp v_\ell$. Because a few residual directions carry a large, slowly drifting share of state variance (the top direction alone is 32% of state variance on the factual store, 96% on generic text), a linear map can predict Δ_ℓ from $\Delta_{\ell-1}$ partly by reading this shared persistent component — running-sum drift, not depth computation. Every AR-on-Δ number in this paper is therefore reported twice: raw, and after the tautology accounting of Section 3.4. The pre-registered scrambled null does *not* separate the two (permuting examples independently per depth destroys genuine structure and tautology together), which is exactly why the pressure battery exists.

3.2 Cheap diagnostic: ridge from the state, against a calibrated lens

The protocol's first, cheap question: is Δ_ℓ linearly predictable from v_ℓ, better than trivial baselines *and* better than what a calibrated per-layer lens already implies? For each layer pair we fit low-rank ridge $\Delta_\ell \approx W v_\ell$ (ranks 16–128) on the train split and score held-out R^2 against: zero-prediction ($R^2 = -1.54$), per-layer mean ($R^2 = 0$ by construction), the identity-transparency regression of $v_{\ell+1}$ on v_ℓ ($R^2 = 0.965$; reported for transparency, dominated by the residual connection, not a comparator), and the **Tuned-Lens-implied update** $\hat{\Delta}_\ell^{\mathrm{TL}} = \hat{r}_{\ell+1} - \hat{r}_\ell$ from the public lens translators (Belrose et al., 2023), the pre-registered gate baseline for proceeding to the expensive AR tests.

Layer-averaged held-out R^2 is **0.602** at the sweep's top rank (128; the untruncated ridge reaches 0.921), versus −0.691 for the Tuned-Lens-implied update (Figure 1). Both are averages of per-layer R^2; the variance-weighted pooled version of the same Tuned-Lens baseline is −0.309 and is the number quoted in Section 6.3. The comparison is not uniform in depth: the lens-implied update is ahead at six mid-stack pairs ($\ell \in \{6\text{–}9, 12, 14\}$, where its per-layer R^2 is positive) and the ridge at the other eighteen, whose catastrophic lens failures dominate the average. The pre-registered gate is the layer-averaged contrast, and it passes with a wide margin — the cheap diagnostic separates cleanly from calibrated per-layer refinement — so the AR programme is justified. A shrinkage-CCA backstop (protocol Step 4) finds held-out leading canonical correlations ≈ 0.9997 at every layer, collapsing to ≈ 0 under the scrambled null; we flag immediately that this near-1 is largely the definitional $-v_\ell \subset \Delta_\ell$ term amplified by the high-variance directions, so CCA confirms shared subspace and exposes nothing ridge missed — the substantive claim rests on the ridge-vs-baselines comparison.

[1] Via the s–r–o-ordered mirror `NeelNanda/counterfact-tracing`; its 21,919 records deduplicate to 21,601 unique (subject, relation, prompt) triples.

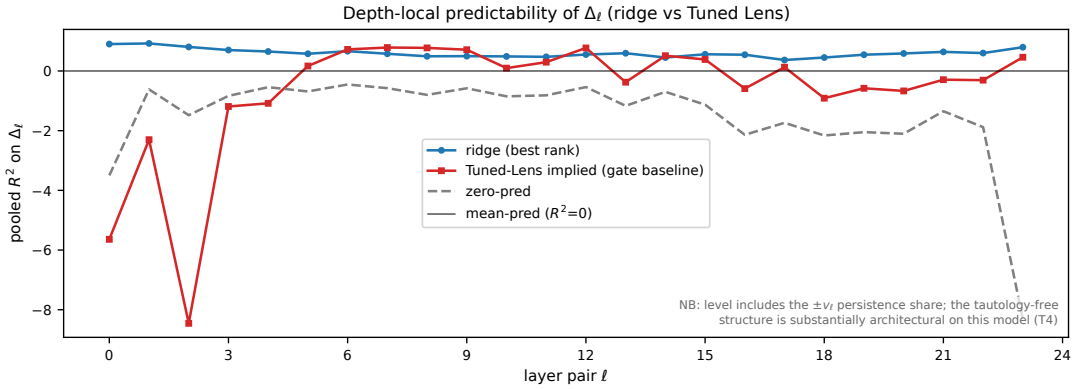


Figure 1: Per-layer-pair held-out R^2 of low-rank ridge predicting Δ_ℓ from v_ℓ on Pythia-410m-deduped, against the mean, zero, and Tuned-Lens-implied baselines. The ridge fit clears the pre-registered (layer-averaged) Tuned-Lens gate with a wide margin (0.602 vs −0.691); per layer the ridge leads at 18 of the 24 pairs, the lens-implied update at six mid-stack ones (Section 3.2). The in-panel note ties the raw level to the tautology accounting of Section 3.4 (Figures 4 and 5), which any reader of this figure should hold alongside it.

3.3 H1 proper: autoregression on Δ and Granger reach

We fit pooled AR(k) models

$$\hat{\Delta}_\ell \;=\; \mu_\ell + \sum_{j=1}^{k} W_j\, \Delta_{\ell-j}, \qquad k \in \{1, 2, 3, 5, 10, 15, 20\},$$

with per-depth intercepts μ_ℓ (so the kernel cannot absorb per-depth mean offsets $\bar{\Delta}_\ell$ — essential for the Phase-2 stationarity comparison) and low-rank W_j shared across depth; predictors are compressed to the top-256 Δ principal components for tractability while the target stays full-D. Held-out R^2 is ≥ 0.46 at every k in the sweep (minimum 0.467 at k=5) and peaks at **0.65** at $k = 20$, rank 128 (Figure 2).

A depth-subset caveat that governs every number below. A pooled AR(k) fit can only score target depths $\ell \geq k$, so the evaluation set *shrinks with* k: at $L = 24$, the k=1 fit averages over 23 depths and the k=20 fit over the four deepest updates (ℓ = 20–23), which happen to be the most predictable. The rise from $R^2 = 0.52$ (k=1) to 0.65 (k=20) is therefore *not* evidence that long lags buy reconstruction. Holding the depth set fixed at $\ell \in [20, 23]$ and re-sweeping k reverses it: AR(1) is the *best* fit there ($R^2 = 0.684$), and every longer lag order is worse (0.647 to 0.655; k=20 gives 0.650). The confounded sweep gains +0.129 from k=1 to k=20; the depth-controlled one *loses* 0.034. This agrees with the per-depth sweeps of Section 4.2 (lags ≥ 2 add a median +0.018 at a single depth) and it makes the "statistically deep, reconstruction-shallow" reading stronger, not weaker. We keep the pre-registered "best-(k, rank)" headline for continuity with the frozen protocol, but from here on we say which depth set a number is scored on, and the k^* verdict rests on the per-depth fits alone.

How far back does depth dependence reach? At every feasible (depth ℓ, lag j) cell we run a nested-F lag test (Wilks' Λ with Rao's F, on a top-32 PCA projection since $D \gg N$ makes full-D Wilks singular): does lag j add signal beyond lags $1, \ldots, j-1$ at depth ℓ? All **112/112** feasible cells are Bonferroni-significant ($\alpha = 4.5 \times 10^{-4}$ per cell), out to the longest lag tested, 20 (Figure 3). Section C verifies the F thresholds are calibrated with a Freedman–Lane permutation check (12/12 spot cells at the permutation floor; parametric and permutation tail quantiles agree within 1%).

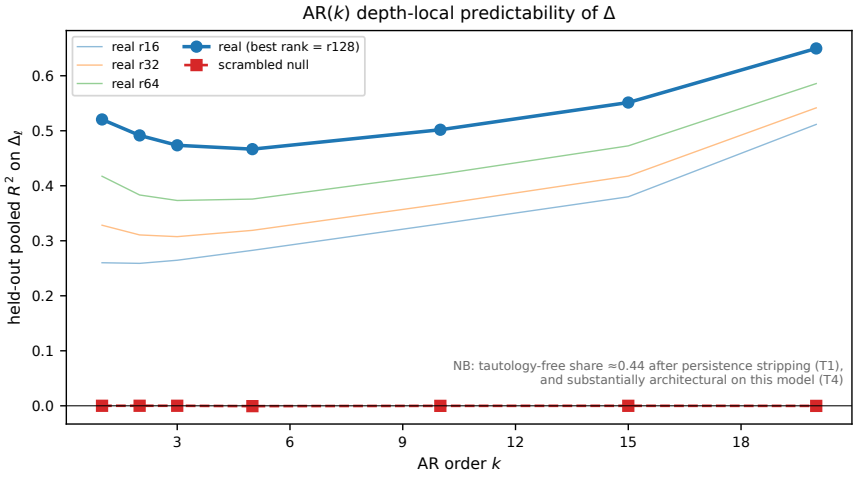


Figure 2: Pooled AR(k)-on-Δ held-out R^2 versus k (real vs the scrambled null) on Pythia-410m-deduped. Real ≥ 0.46 everywhere, peaking at 0.65; the null is indistinguishable from zero. Two corrections travel with this panel. Each k is scored on the depths $\ell \geq k$, so the curve is not a lag effect: fixing the depth set at $\ell \in [20, 23]$ makes k=1 the *best* fit (0.684), with longer lags costing up to 0.037 (Section 3.3). And the raw level splits into tautological and genuine parts (Figures 4 and 5), the genuine part being substantially architectural on this model.

Null 1 — scrambled activations (pre-registered). Permuting the example axis independently at each depth preserves every per-layer marginal and destroys within-example depth structure. The AR fit collapses to $R^2 \approx 0.0000$ at every k and the Granger grid to **0/112** significant cells. The signal is example-specific cross-depth structure, not a marginal artefact.

Null 2 — context perturbation (pre-registered, the critical control). Each prompt is wrapped in a fixed six-token neutral preamble; the perturbed variant replaces the preamble's content words with random function words at identical absolute positions. On the 863/1000 sampled examples where the wrapped baseline still recovers the answer, the AR fit does not degrade (k=20: 0.639 wrapped $\rightarrow$ 0.696 perturbed; k=1, the whole-stack fit: 0.515 $\rightarrow$ 0.496) and the Granger heatmaps correlate at **0.964** (112/112 $\rightarrow$ 109/112 significant). The cross-layer dynamics at the query token are a property of subject–relation processing — depth-intrinsic, not leaked context.

Two honesty notes, because this is the control the whole programme leans on. First, these are *raw* AR levels, and Section 3.4 shows raw levels are dominated by the running-sum tautology; the control's force comes from the Granger heatmap correlation and the lag-$\geq$2 grid (both tautology-light), not from the R^2 agreement. Second, the k=20 level *rises* under perturbation. On this store that is the persistence signature — less informative context makes the residual stream drift more smoothly, and Section 3.4 shows generic text scores *higher* raw (0.935) while carrying ≈ 0 genuine structure. We therefore read Null 2 as "the depth structure is not inherited from the context", and not as "perturbation improves the model".

3.4 Pressure-testing H1: how much is tautology, and is the rest learned?

Gate 1 passed on the pre-registered criteria. Before building on it, we attacked the headline with five post-hoc controls (T1–T5): four tautology accountings and an untrained-network control (T4). The corrected reading below, not the raw 0.65, is the version of H1 the rest of the paper uses.

T5 — exact tautology floor. The held-out explained variance of the single headline AR fit (k=20, so scored on $\ell \in [20, 23]$) splits exactly (orthogonally, Pythagorean) across the top-m

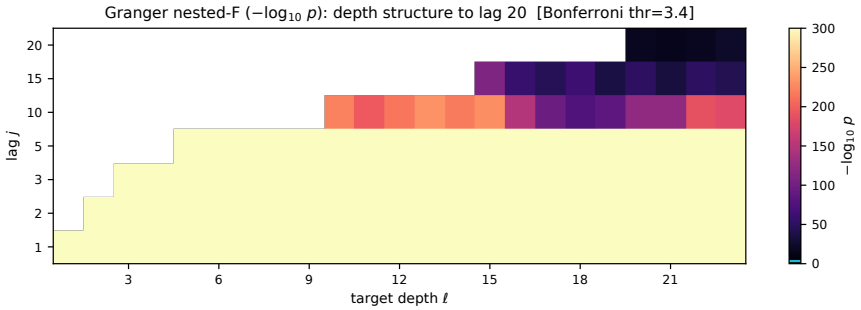


Figure 3: Granger-style nested-F significance ($-\log_{10} p$) per (depth, lag) cell. All 112 feasible cells are Bonferroni-significant out to lag 20; the scrambled null yields 0/112 (not shown). The Bonferroni cut ($-\log_{10} p \approx 3.4$) is marked on the colour bar: the deepest lag-20 cells render dark on this 0–300 scale yet still clear it. Statistically the depth dependence is long-range — reconstruction-wise it is first-order (Section 4.2), and Section 3.5 localises the long reach to the attention write.

persistent state subspace and its complement. At $m = 50$: **factual** 0.65 = 0.52 **tautological** + 0.13 **genuine**; the persistent subspace itself is near-trivially predictable ($R^2_{\|} = 0.90$–0.95 over the m sweep), which *is* the $\pm v_\ell$ drift made explicit (Figure 5). The four deepest updates are where the stream drifts most, so this is the tautology at its worst; repeating the decomposition on the whole-stack k=1 fit gives 0.52 = 0.27 tautological + 0.25 genuine (Table 3). The tautological share is the clear majority on the headline fit; on the whole stack it is a bare majority (0.272 vs 0.248, a margin that carries no interval), which we round to "roughly half". We quote both because the pre-registered headline is the four-depth number.

T1 — persistence-stripped refit. Projecting out the top-m principal directions of the *states* and refitting: removing the single top direction (32% of state variance) costs only $0.650 \to 0.637$; removing the top 50 (87% of state variance) still leaves $R^2 = 0.436$, roughly 4000× this sweep's own scramble floor ($\sim 1 \times 10^{-4}$ on the primary factual store; the floor is store-dependent, and Table 3 reports the largest we saw, 7×10^{-3} on GPT-2) (Figure 4). Two reasons T1 exceeds T5's genuine share: it re-optimises the kernel on the complement, and — since each point is a max over the k sweep — stripping flips the selected order from k=20 to k=1, i.e. onto the 23-depth evaluation set. Section C therefore repeats the whole m sweep at fixed k=1 for every store used below, and all trained-vs-untrained and factual-vs-generic contrasts we quote are taken from that like-for-like table.

T2 — persistence-matched surrogate. A random walk with independent per-layer Gaussian steps matched to each Δ_ℓ's mean and (rank-256) covariance — all persistence, no cross-layer dependence — scores $R^2 = \mathbf{0.0000}$. State persistence alone reproduces *none* of the AR-on-Δ predictability (for AR on increments, the $\mathrm{Var}(v_\ell)$ contributions to $\mathrm{Cov}(\Delta_\ell, \Delta_{\ell-1})$ cancel in expectation under independent steps), so the scramble floor was already the honest floor.

T3 — lag≥2 Granger, free of the shared term by construction. The shared $\mp v_\ell$ term couples only lag 1 ($\Delta_{\ell-2}$ shares no state term with Δ_ℓ). Restricting the grid to lags ≥ 2: **89/89** cells significant (real) versus **0/89** (null), identically on factual and generic stores. Depth dependence beyond the mechanical overlap is total. (Strictly, T3 removes the algebraic overlap, not the correlation that persistence induces between Δ_ℓ and $\Delta_{\ell-2}$; T2 is what rules the latter out, by showing that persistence with no cross-layer dependence scores zero. The two controls are complements.)

T4 — untrained-network control (deflationary, and one-family). Running the identical battery on a *randomly initialised* Pythia-410m-deduped (same prompts, same pipeline; a single

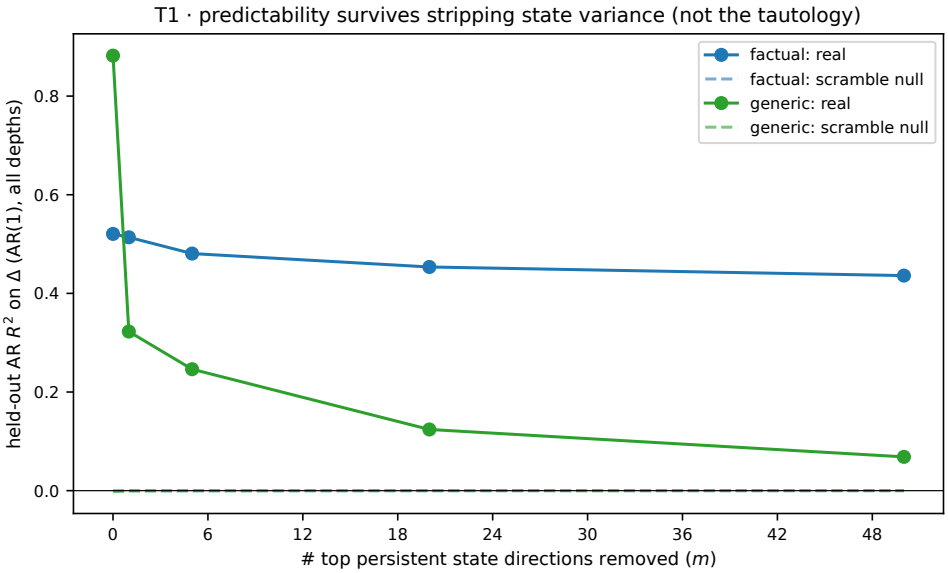


Figure 4: T1 persistence-stripped AR R^2 as the top-m state directions are removed, factual vs generic stores. Both curves are AR(1) fits over all 23 feasible depths, so they are directly comparable (the max-over-k version used elsewhere is not: it selects different depth sets for the two stores; Section C). The factual curve decays gradually and saturates at 0.44 ≫ null; the generic store — whose top direction alone is 96% of state variance — collapses on the very first direction and reaches 0.07 by m=50. Nearly all of the raw generic level was persistence, and what survives is an order of magnitude weaker than the factual store's.

initialisation, Section 7), at a common k=1 so that trained and untrained are scored on the same 23 depths: the persistence-stripped (m=50) AR R^2 is **0.503 (untrained) vs 0.436 (trained)** on factual prompts, and the untrained network passes every tautology-free control — scramble ≈ 0, surrogate ≈ 0, lag≥2 Granger 89/89 — exactly like the trained one. The pre-registered inference rule ("trained ≫ untrained ⇒ learned") returns the other branch: on the primary model, **the tautology-free AR-on-Δ structure is substantially architectural** — a property of the residual stack with per-layer nonlinear maps and LayerNorm, not of the learned weights. What training *does* change is the persistent/drift component (raw R^2 0.543→0.521 factual, where the tautological share rises 0.12→0.27; 0.063→0.882 generic, almost all of it tautology) — the part T1/T5 already attribute to drift.

It does not replicate on GPT-2-small, and the earlier claim that it did was an artefact of the depth-subset confound. Compared at the same k=1 over the same 11 depths, GPT-2's persistence-stripped R^2 is **0.439 trained vs 0.233 untrained** — a learned share of 47%, not the ∼5% we obtain if the untrained arm is allowed to select k=10 and thereby score only its two deepest updates. So the honest statement is narrower than we first drew it: *on Pythia-410m-deduped the tautology-free AR structure is architectural (the untrained network exceeds the trained one); on GPT-2-small it is about half learned.* We have no account of why the two families differ, and we flag it as the most interesting loose end in Section 3. Table 3 carries every number.

3.5 Localising the depth reach: attention vs MLP writes

Splitting each update into its two additive sublayer writes, $\Delta_\ell = \mathrm{attn}_\ell + \mathrm{mlp}_\ell$ (exact for Pythia's parallel blocks; additivity holds to 7×10^{-8} relative error), gives streams that carry *no* $\pm v_\ell$ term. Both independently reproduce H1 (AR real−null on the four deepest updates: attention **0.542**, MLP **0.653**; Granger 112/112 vs 0/112 both), and they dissociate cleanly: the **MLP write is strictly first-order at every depth** ($k^* = 1$) while the **attention write carries the long-range depth memory** (k^* median 2, rising to 10 at ℓ=22), with MLP contributing 78%

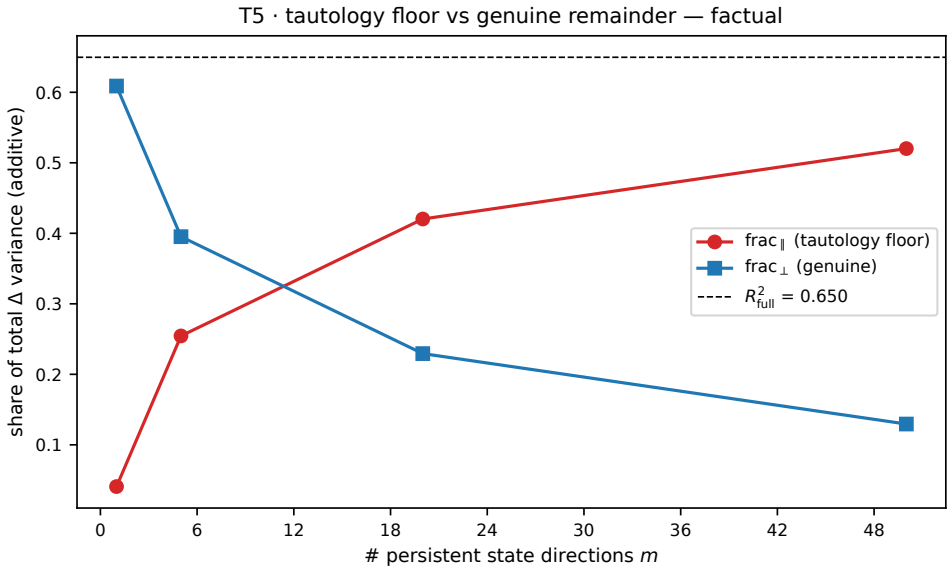


Figure 5: T5 exact decomposition of the single headline AR fit (k=20, scored on $\ell \in [20, 23]$) into tautological ($R^2_{\parallel}$, persistent-subspace) and genuine ($R^2_{\perp}$) shares as a function of the subspace size m. Factual (shown): $0.65 = 0.52 + 0.13$ at m=50. Generic (not shown): $0.935 = 0.937 - 0.002$, a marginally *negative* genuine share — the generic level is persistence, not computation. At k=1, over all 23 depths, the same split gives $0.521 = 0.272 + 0.248$ (factual) and $0.882 = 0.879 + 0.003$ (generic); Table 3.

of update energy. The residual stream's statistically-deep-but-reconstruction-shallow tension (Figure 3 vs Section 4.2) is therefore the energy-dominant memoryless MLP channel superposed on a small deep-reaching attention channel.

What this split does and does not establish. It removes the *mechanical* $\pm v_\ell$ overlap, and that is all. In Pythia's parallel block both writes are deterministic functions of the same state, $\text{attn}_\ell = f_\ell(\text{LN}(v_\ell))$ and $\text{mlp}_\ell = g_\ell(\text{LN}(v_\ell))$, and v_ℓ is strongly persistent, so consecutive sublayer writes remain correlated *through* that shared, slowly drifting argument. The split is therefore a tautology-free *localisation* of where the depth reach lives, not an independent demonstration that the AR structure is persistence-free. The persistence-free evidence remains T1–T3 below, and T4 shows the surviving structure is largely architectural — exactly what a smooth per-layer map of a persistent state would produce. We did not run the T2 surrogate or the untrained control on the sublayer streams; doing so is the natural next check.

3.6 Robustness and the Gate-1 verdict

An object-disjoint split (grouping by answer object, since objects like "French" recur across up to 163 prompts while subjects are nearly unique) leaves the result intact: AR $R^2 = 0.561$, Granger 112/112, null ≈ 0. H1 also replicates on all five replication models (Section 4.4) and on length-matched generic Pile text (Section 4.5).

Gate-1 verdict (pre-registered routing): H1 supported; proceed to Phase 2 — with the corrected interpretation carried forward: genuine, above-null, depth-intrinsic cross-layer structure exists in the updates (T1–T3); on the primary model it is substantially a property of the architecture rather than of training, though that reading is Pythia-specific (T4); and raw AR-on-Δ levels must never be quoted without the tautology split (T5) or without the depth set they were scored on (Section 3.3).

4 Phase 2 — what kind of structure? Stationarity, order, and scope

4.1 H2 fails: the kernel is pervasively depth-dependent

H2 asks whether the AR kernel is the *same* kernel at every depth. The pre-registered test partitions depth into thirds (early ℓ 0–7, middle 8–15, late 16–23), fits one band-shared AR(k) kernel per band (per-depth intercepts μ_ℓ absorb mean offsets, and are *refit from the train split at the evaluation depths* whenever a kernel is transferred, so only transition structure is compared; $k = 5$, the plan's early-band feasibility cap $\lfloor L/3 \rfloor - 3$, which leaves the early band three scorable target depths, $\ell = 5$–7; working rank 256 after the plan's r-vs-$2r$ spot-check promoted 128), and compares the kernels four ways.

Every comparison rejects stationarity (Figure 6): *cross-band prediction* — within-band held-out R^2 is positive (0.74 / 0.58 / 0.67) but **every** off-diagonal is **negative** (e.g. early→middle -1.00, late→middle -1.89): a kernel from one band predicts another band's updates *worse than that band's own mean.* This holds at every truncation rank from 32 to 256; the statistic that carries the claim is the *least* negative off-diagonal, which is still -0.15 at rank 32 (rank truncation shrinks every cell toward zero, so low rank is the hard case). *Relative Frobenius distance* $\|W_a - W_b\|_F / \frac{1}{2}(\|W_a\|_F + \|W_b\|_F)$ between band kernels is 1.39–1.42, straddling the value $\sqrt{2}$ that two orthogonal equal-norm kernels would take (the metric's ceiling is 2, at $W_a = -W_b$); *principal angles* between kernel output subspaces average 57–62° — well short of the 90° two random subspaces would give, so the kernels share structure even as their transfer collapses; and a *bootstrap* within bands puts the between-band distance **80–104 standard deviations** above within-band sampling spread. Crucially there is *no clean phase boundary*: the three pairwise distances are equal to within 3%, so adjacent bands differ as much as distant ones — the depth dependence is **pervasive drift**, not segmentation.

Three caveats travel with this result: the cap makes it a short-horizon test (long-horizon stationarity is rechecked, and also fails, on the fitted generative model); AR-on-Δ stationarity is a proxy for H2's generative (MA) form — but a proxy whose failure is decisive in the rejecting direction; and the $\pm v_\ell$ shared term biases the comparison *toward* stationarity, so rejection is the conservative outcome.

Gate-2 verdict (pre-registered decision matrix): H2 fails pervasively ⇒ Phase 3 fits the matrix's state-space branch — a state-space model with *free per-layer operators* (the smoothly-varying variant assumes exactly the smoothness the data reject), and — per the proposal's own framing — CLT-style unconstrained per-(source, target) parameterisation is *validated*: "shared kernel along depth" is the wrong abstraction for these dynamics. The parsimony claim originally attached to H2 is off; what survives is a reconstruction comparison at stated parameter budgets (Section 6.4).

4.2 Effective order: $k^* \approx 1$, everywhere we looked

Per-depth AR(k) sweeps — one fit per target depth, so no depth-subset confound (Section 3.3) — define the effective depth k^* as the smallest k within 95% of the asymptotic R^2. On the factual store $\boldsymbol{k^* \approx 1}$ at every depth (median 1; the single exception, k^*=2 at ℓ=17, coincides with the answer-promotion write discussed in Section 6.7): one preceding update captures essentially all held-out predictable variance (per-depth best R^2 0.83–0.94 away from ℓ=17, where it falls to 0.69), lags ≥ 2 add a median +0.018 (maximum +0.050, at ℓ=17), and the scrambled null collapses the whole curve to 0 (Figure 7). This per-depth measurement, not the pooled sweep of Section 3.3, is what licenses the contrast with the Granger reach of 20 (Figure 3): **long-range depth structure is statistically real but reconstruction-negligible** — and Section 3.5 locates the two halves of that sentence in different sublayers (deep reach: attention; reconstruction mass: MLP). k^* inherits the lag-1 tautology caveat, but it is also exactly the model-order information Phase 3

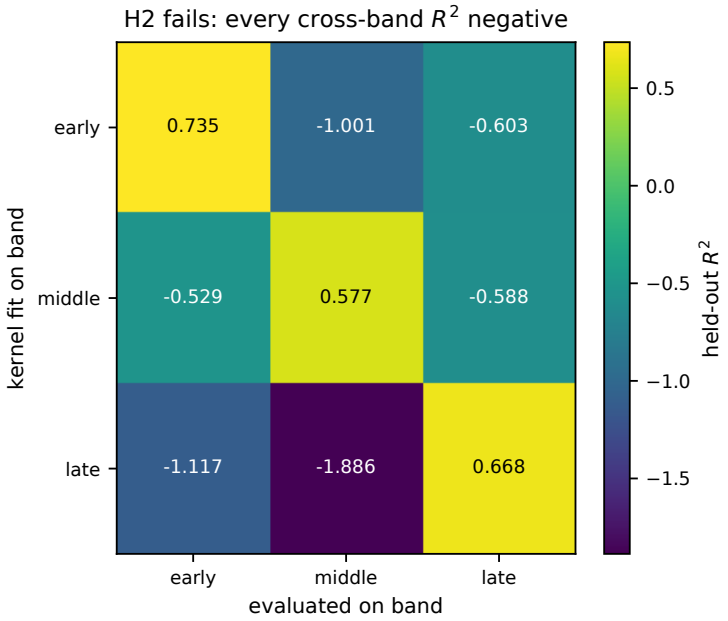


Figure 6: H2 stationarity test on Pythia-410m-deduped: held-out R^2 of band-shared AR(5) kernels evaluated within and across depth bands. Positive diagonal, every off-diagonal negative — a kernel from one depth band is worse than the target band's mean. Rank-robust; bands bootstrap-separated by 80–104 SD; no segment boundary.

needs, and the fitted model later supplies a candidate mechanism for it (a fast-attenuating state path, Section 6.8). Stratifying by relation class leaves $k^* = 1$ in every class and every depth band (class-invariant order); on generic text the per-depth median is $k^* = 2$ (pooled band-level fits give 1/3/1, and lags ≥ 2 add a median +0.039) — stated explicitly since it differs from the factual 1, with per-depth generic fits being visibly noisier.

4.3 Scope: the kernel is query-position-specific and circuit-specific

Token position. Re-collecting trajectories at four positions per prompt (first subject token, last subject token, a mid-prompt non-subject token, final prompt token) and refitting AR(k^*=1) per position: within-position held-out R^2 is positive at every position (0.28–0.55), while **eleven of the twelve cross-position transfers are negative** (mean −0.36; the lone exception, last-subject-token → mid-prompt, reaches +0.024, i.e. it recovers 5% of that position's own R^2). The kernels separate exactly as the depth bands do — relative Frobenius 1.25–1.46, straddling the $\sqrt{2}$ of orthogonality, with mean principal angles 57–60° that (as in Section 4.1) stay well inside the 90° of independent subspaces: shared structure, no transfer, now across positions. The scrambled null collapses the diagonals, and the tautology, being shared by the diagonal and the off-diagonal alike, is discounted — though, entering additively, not exactly cancelled — by the *retention ratio* (cross-position R^2 divided by the target position's within-position R^2; worst −2.88). The Phase-1/2 characterisation is therefore scoped to *context-aggregated factual recall at the query position*; Phase 3 fits there, and every claim in this paper inherits that scope.

Relation class. The mirror-image test across the three relation classes (run with the fitted SSM of Section 5; reported here because it completes the scope picture): within-class held-out R^2 mean +0.835, **every cross-class transfer ≈ 0 or negative** (mean −0.191; Figure 8). The dynamics are finely tuned to each recall circuit — *circuit-specific kernels* — while the *order* $k^* = 1$ is class-invariant: same model order everywhere, different kernels per circuit.

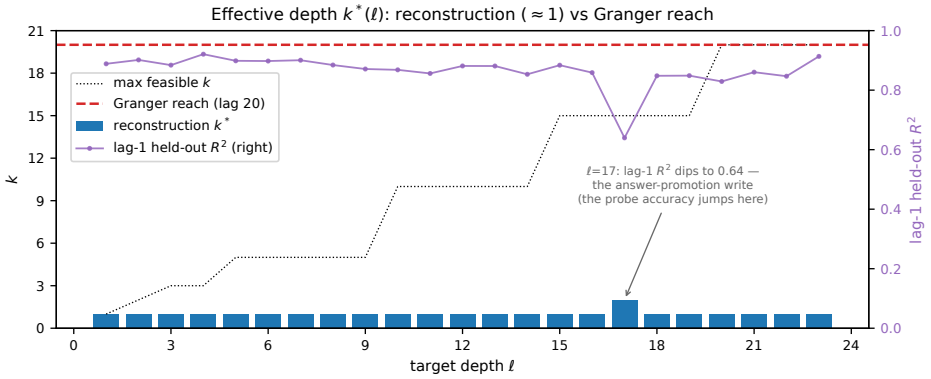


Figure 7: Per-depth effective order k^* on Pythia-410m-deduped (bars), with the per-depth lag-1 R^2 (purple line, right axis) and the Granger reach (20) for contrast. The one k^*=2 depth, ℓ=17, is the answer-promotion write: the state probe jumps there (Figure 14) and the SSM's per-layer reconstruction R^2 dips there (Figure 16, right panel).

Table 1: Cross-model replication of H1/H2. AR R^2 is pooled held-out real−null at the sweep's best lag order k^{sel}, which is reported alongside because a pooled AR(k) fit scores only depths $\ell \geq k$ (Section 3.3): the levels are comparable *within* a row against its null, not *across* rows with different k^{sel}. What replicates is the sign pattern, not the level. Granger counts are Bonferroni-significant cells (real vs scrambled null) out of all feasible (depth, lag) cells — $\sum_j (L-j)$ over the lag grid: $\{1,2,3,5,10,15,20\}$ for $L \geq 24$ (112 cells at L=24, 196 at L=36); the 12-layer models use the capped grid $\{1,2,3,5,8,10\}$ — lags 15 and 20 are infeasible at L=12, and lag 8 is added to keep a six-point grid — giving 43 cells (Section C). The cross-band column reports the sign pattern of the H2 transfer test. Datasets are per-model confidence-filtered CounterFact (sizes 767–2,540). The tautology accounting of Section 3.4 was run on two of these models (Pythia-410m-deduped, GPT-2-small), not all six; these AR levels are raw.

Model	L	AR R^2 (real−null)	k^{sel}	Granger real / null	cross-band R^2	k^*
GPT-2-small	12	0.625	1	43/43 / 0	all < 0	1
Pythia-160m-deduped	12	0.894	10	43/43 / 0	all < 0	1
GPT-2-medium	24	0.513	20	112/112 / 0	all < 0	1
Qwen2-0.5B	24	0.462	1	112/112 / 0	all < 0	1
Pythia-410m-deduped (primary)	24	0.650	20	112/112 / 0	all < 0	1
GPT-2-large	36	0.311	1	196/196 / 0	all < 0	1

4.4 Cross-model replication: six models, three families

The full H1/H2 pipeline (and, where public residual SAEs exist, the Phase-3 fit) was replicated on five further models; Pythia-70m additionally serves as the optimiser-debug lane and a depth point in Section 6.2. Table 1 shows the pattern is universal across the models tested: AR-on-Δ predictability far above its null with full Granger significance, every cross-band R^2 negative, and $k^* = 1$ at every band of every model.

One anomaly is recorded rather than hidden: on GPT-2-small the cheap *low-rank* ridge diagnostic scores only 0.146 layer-averaged despite a full-rank 0.91 and AR 0.625 — GPT-2's $v \to \Delta$ map is far less low-rank than Pythia's. It does not affect any H1/H2 verdict.

4.5 Generality beyond factual recall

Re-running the identical pipelines, on the primary model, on 1,555 generic Pile windows — length-matched to the factual prompt multiset (content lengths 4–24 tokens, median 9; one window per distinct document; this narrow-window construction is a stated scope condition) — replicates the *entire* characterisation: H1's binary claim holds (null ≈ 0, Granger 112/112 vs 0/112, lag$\geq$2 89/89), H2 fails even more emphatically (every cross-band R^2 negative, worst -144; bands separated by 39–188 bootstrap SD), and the effective order is low (k^* median 2).

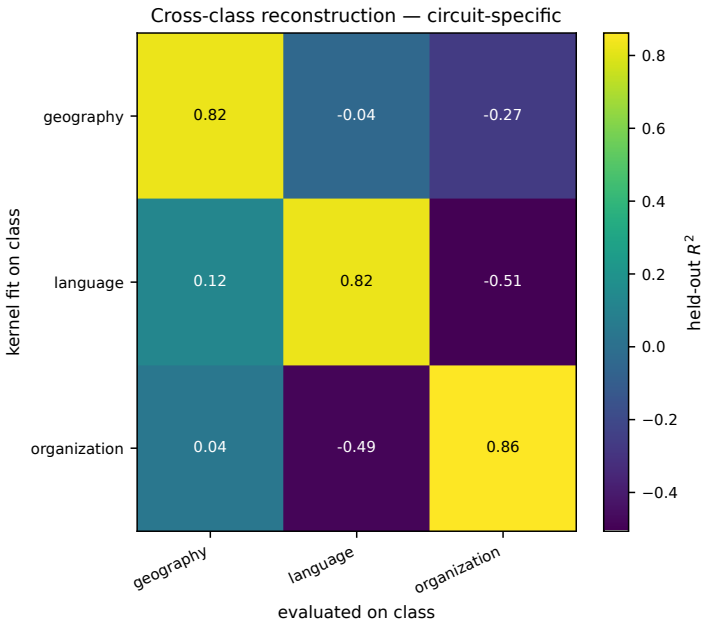


Figure 8: Cross-class transfer of the fitted decomposition (train class → eval class, held-out R^2). Strong diagonal, at most weak transfer off it (largest off-diagonal +0.12, 15% of that column's within-class value) — the same signature as cross-band and cross-position: structured, but specific.

The raw generic AR level (0.935) is *higher* than factual (0.650), and Section 3.4 is the reason that must not be read as "more computation": T5 shows the generic level is ≈ 0.94 tautological + ≈0 genuine, and at a common lag order the persistence-stripped generic structure (0.07) is an order of magnitude *weaker* than factual (0.44). So what generalises is H1's binary claim — above-null, example-specific, depth-intrinsic structure exists in the updates of generic text too — and emphatically not its magnitude, which on generic text is almost entirely running-sum drift. The phenomenon is general to transformer depth; the *kernels* remain position- and circuit-specific per Section 4.3.

5 The decomposition: a first-order, fully depth-indexed SSM over sparse sources

5.1 Form

Gate 2's decision matrix, applied to Section 4, fixes the architecture: pervasive (not smooth, not segmented) depth dependence rules out shared and interpolated kernels, and $k^* \approx 1$ makes the depth latent first-order. The generative model is a linear state-space model with free per-layer operators:

$$z_{\ell+1} = A_\ell z_\ell + B_\ell u_\ell, \qquad \hat{\Delta}_\ell = b_\ell + C_\ell z_\ell + D_\ell u_\ell, \tag{1}$$

where $u_\ell \in \mathbb{R}^{M_\ell}$ are sparse per-layer source codes, $z_\ell \in \mathbb{R}^{d_z}$ is the cross-layer state, b_ℓ is an unpenalised per-layer bias absorbing $\bar{\Delta}_\ell$, and every operator is rank-factored ($A_\ell = A_\ell^{(1)} A_\ell^{(2)\top}$, rank r; likewise B, C, D). The feedthrough $D_\ell u_\ell$ is the per-layer (no-propagation) path; the state path $C_\ell z_\ell$ carries whatever earlier layers contribute. The comparison that isolates the value of cross-layer propagation is therefore built in: the same model with the state ablated ($\hat{\Delta}_\ell = b_\ell + D_\ell u_\ell$, `use_state=False`) is a matched per-layer sparse dictionary sharing the optimiser, sources, sparsity sweep, and rank. Primary setting: $r = d_z = 32$.

5.2 Sources

Source codes initialise from public per-layer SAE latents of v_ℓ. On the primary model these are the 24 *independent single-layer* TopK SAEs released alongside the MLSAE work for Pythia-410m-deduped (Lawson et al., 2025) — one checkpoint per block, not the paper's shared multi-layer dictionary ($M = 65{,}536$, TopK k=32, all 24 layers); on the narrow query-position distribution only 497–4,433 latents per layer ever fire, so each layer is pruned to its alive set ($M_\ell^* \leq 4{,}433$) — standard dead-latent removal, exact for reconstruction, and legitimate here because the depth-indexed form permits per-layer widths. v_0 has no public embedding SAE, so u_0 is zero-initialised and free. Matched SAE state reconstruction averages $R^2 = 0.598$ across layers (two weak public checkpoints at v_{18}/v_{19}, 0.21/0.31, are shared by both comparison arms and cancel in every gap we report). Because the dictionary turned out to matter for identifiability (Section 6.5), we also train matched **L1** and **TopK** dictionaries on the same cached activations — identical architecture (per-token standardisation, tied decoder bias, unit-norm decoder columns), differing *only* in the sparsity mechanism (L1 coefficient 0.5 giving $L_0 \approx 36$, matched to TopK k=32) — including leakage-free variants fit on the train split only.

5.3 Optimisation and evaluation protocols

The loss is squared reconstruction error over depths and examples, plus an L1 penalty on all $u_\ell^{(i)}$ and a Frobenius penalty on the operators ($\lambda_{\text{reg}} = 10^{-4}$). Optimisation alternates exactly as pre-registered: (i) a proximal-gradient LASSO pass over each example's stacked codes $(u_0, \ldots, u_{L-1})^{(i)}$ with operators frozen — FISTA with backtracking (Beck and Teboulle, 2009), validated against `sklearn` LASSO and closed-form least squares, keeping the exact zeros a penalised Adam would lose — and (ii) a full-batch Adam (Kingma and Ba, 2015) sweep over $\{A_\ell, B_\ell, C_\ell, D_\ell, b_\ell, z_0\}$ with codes frozen, differentiated through the unrolled SSM. Ten outer iterations; the method was debugged end-to-end on Pythia-70m-deduped before touching the primary lane.

Held-out numbers in this paper come from two protocols, and conflating them creates apparent contradictions, so we name them here and tabulate every operating point in Section B:

LASSO protocol (the Gate-3 protocol). Operators frozen after training; test codes re-inferred by the same LASSO against the test Δ. This is *reconstruction at matched sparsity* — the standard dictionary-learning evaluation — not out-of-sample forecasting; both comparison arms are treated identically.

Frozen-source protocol. Decoders fit by Adam on the *fixed* cached codes, no re-inference. Used for the CLT family (Section 6.4), where it isolates the decoder parameterisation; its absolute levels sit ∼0.06–0.08 below the LASSO protocol for every method (Table 2).

5.4 Pre-registered sanity checks

All mandatory checks pass on the primary fit: no collapse to the bias solution (operator norms 252 vs bias norm 11; code density $L_0 \approx 0.009$ as a fraction of alive atoms on the train fit — held-out re-inference at the knee runs denser, $L_0 = 0.0149$ (Section 6.1), and source-SAE L_0s are quoted as per-token counts); per-example reconstruction spot-check median $R^2 = 0.98$; top-20 features have mean relation-class purity 0.83 with human-legible themes (e.g. an "employed-by/manufactured-by → company" feature at layer 9, a "premiered-on → TV network" feature at layer 1, a "European place → Europe" feature at layer 22); dataset stability across disjoint halves top-100 Jaccard 0.89. Seed stability of decoder *directions* failed its pre-registered bar on the TopK source ($0.45 < 0.7$) — recorded at Gate 3 as a limitation, and subsequently traced to the dictionary type and closed, a story told in full in Section 6.5.

Gate-3 verdict: GO — the reconstruction clause is met with margin (next section), and the then-open seed-stability limitation was scoped to single-feature decoder identifiability, which no headline claim uses.

6 Evaluation

Unless stated otherwise, numbers are held-out R^2 on the fixed 311-example test split, and the reconstruction gaps carry a paired bootstrap 95% confidence interval over test examples (10,000 resamples; Section C). Two exceptions, flagged where they occur: the depth series (Section 6.2) reports mean±sd over three refits rather than a bootstrap CI, and the causal group effects (Section 6.6) bootstrap over the eight candidate features, which are the independent unit there. Several supporting numbers — the T1–T5 battery, the sublayer split, the Qwen full-depth contrast — carry no interval at all. Model selection performed on the test split is limited to points where the contrast itself is the pre-registered quantity; the one CI where selection could matter (SSM vs CLT) is re-derived with validation-based selection in Section 6.4. Section B tabulates every operating point and which claim uses which.

6.1 Headline: cross-layer state beats a matched per-layer dictionary

Under the LASSO protocol, the SSM sits above its state-ablated twin at every sparsity level of the λ_{L1} sweep (Figure 9): held-out R^2 0.867–0.698 versus 0.816–0.653 across $\lambda_{L1} \in \{0.3, \ldots, 30\}$. At the canonical λ_{L1}=1 knee[2] the gap is

$$\textbf{+0.0489}\ [+0.0464,\ +0.0516]\ \text{(MLSAE source)},$$
$$\textbf{+0.0483}\ [+0.0458,\ +0.0509]\ \text{(trained-L1 source)},$$

— dictionary-independent, seed-robust (three seeds give $+0.0504 \pm 0.0014$; the CIs above are seed 0, the recorded canonical fit), and it survives an object-disjoint split (+0.055/+0.061 matched-L_0 at the two central sparsities). At the knee the two arms are not exactly matched in sparsity (SSM L_0 0.0149 vs baseline 0.0133 as a fraction of alive atoms); interpolating the baseline Pareto to the SSM's L_0 moves the gap to +0.048, so the mismatch is not doing the work. The same interpolation across the full sweep spans +0.037 to +0.066, but its maximum is upward-biased (it pairs against a baseline tail that has saturated in L_0), which is why the knee gap, not that range, is the statistic we carry forward (Section A, deviation 5).

Matched parameters, not just matched rank. The gap above compares arms at equal rank, giving the SSM roughly twice the baseline's parameters (8.36M vs 4.22M). Widening the baseline to rank 64 (8.41M) closes most of it: under the LASSO protocol the state-path edge survives at +0.0077 [+0.0053**,** +0.0102] — small, real, and conservative (the wider baseline runs slightly denser) — while under the frozen-source protocol the same comparison is a wash (±0.005, opposite signs at the 8.4M and 16.8M budgets). We therefore quote the headline as "+0.049 **rank-matched,** +0.008 **param-matched**" and never write "load-bearing" without the second number. What decoder width cannot replicate at any budget is the *structure* the state provides — Sections 6.7 and 6.8.

6.2 The gain grows with depth

Repeating the fit at every depth in the model zoo and reporting the *knee-gap* statistic (the raw same-λ gap at the $\lambda_{L1} = 1$ knee; the earlier max-over-λ statistic is upward-biased and seed-fragile because the baseline Pareto's tail saturates in L_0 — both statistics are archived) gives a monotone in-family depth series with tight 3-seed error bars:

$$\text{Pythia: } 6L\ \ +0.0076 \pm 0.0012\ \rightarrow\ 12L\ \ +0.0130 \pm 0.0019\ \rightarrow\ 24L\ \ +0.0504 \pm 0.0014,$$

[2] "Knee" is the sweep's sparsity–reconstruction bend, visible in Figure 9: relative to λ_{L1}=0.3, λ_{L1}=1 buys a ∼3× lower L_0 for −0.010 R^2, while beyond it L_0 saturates and R^2 only falls. The same bend appears on the Pythia-70m debug lane, and λ_{L1}=1 was frozen as the canonical operating point when the Phase-4 checkpoint was persisted — before the knee gap became a reported statistic — so the operating point was not selected on the gap it now reports.

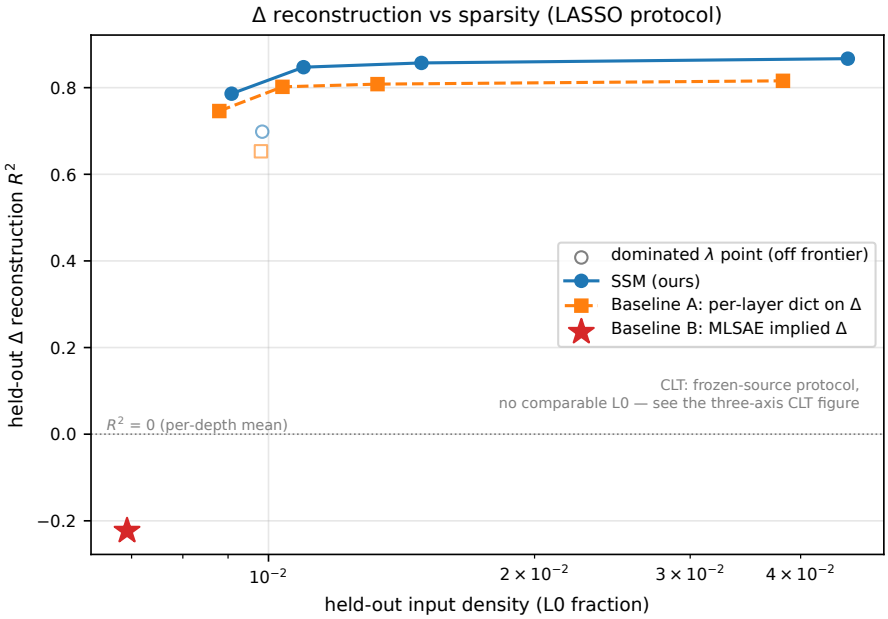


Figure 9: Sparsity–reconstruction Pareto on Pythia-410m-deduped (LASSO protocol): the SSM vs its state-ablated per-layer twin (Baseline A), with the naive differenced-SAE baseline (Baseline B, negative) for context (Section 6.3); the CLT lives in the frozen-source protocol with no comparable L_0 and gets its own figure (Figure 11). The SSM is above the matched baseline at every operating point.

with cross-family 12-layer points GPT-2-small $+0.0175 \pm 0.0025$ and Qwen2 $+0.0413 \pm 0.0025$ (Figure 10). These are single points, one per family: they replicate the gain's *sign*, not its growth with depth, which is a within-Pythia statement from three points.

Five caveats are part of this figure's claim, and together they are the reason it is offered as suggestive rather than settled. (i) The Pythia line mixes two TopK suites (EleutherAI-32k at 70m, MLSAE elsewhere), and the cross-family points use L1 suites, so dictionary type is not held fixed anywhere in the figure. (ii) Within a family, depth is confounded with width and parameter count (deconfounding via Pythia-1b/1.4b is recorded future work). (iii) Training-set size also varies along the Pythia line ($N = 1{,}000 \to 767 \to 1{,}555$ at 6/12/24 layers), and Section 6.4 shows the relative value of cross-layer parameterisations moving strongly with N in exactly this range (the SSM−CLT gap flips sign); we did not measure the knee gap's own N-dependence, so N stays a live confound for the depth trend. (iv) The 6-layer point is the Pythia-70m debug lane, whose dataset is built differently (top-1,000 candidates by confidence rather than threshold-filtered; Section D). (v) The plotted Qwen point is the 12-layer *SAE-coverage prefix of a 24-layer model* — a dictionary-controlled contrast (our own L1 SAEs on the same rows at 24 vs 12 layers) gives $+0.072$ **at full depth vs** $+0.043$ **on the prefix**, so public-suite coverage, not depth, limits the plotted point. That $+0.072$ is an L1-sourced number and Pythia-410m-deduped's $+0.050$ is TopK-sourced, so we do not read the two against each other.

How much of the gain is learned? Running the identical pipeline on a *randomly initialised* Pythia-410m-deduped (sources retrained on its own activations) yields a positive but smaller gap: mean over the λ sweep $+0.025$ **untrained vs** $+0.046$ **trained** — the untrained network carries **53%** of the trained gain (per-λ shares averaged), and the learned share *grows with sparsity* (66% learned at λ_{L1}=10). The split replicates on GPT-2-small (untrained carries 44%). Read against Section 3.4, this is a dissociation *on the primary model only*: there, AR-on-Δ *predictability* is architectural (the untrained network *exceeds* the trained one once persistence is stripped) while the SSM *reconstruction win* keeps a ≈half learned component — so the decomposition sees something training adds that the raw AR statistic does not, concentrated at the sparse operating

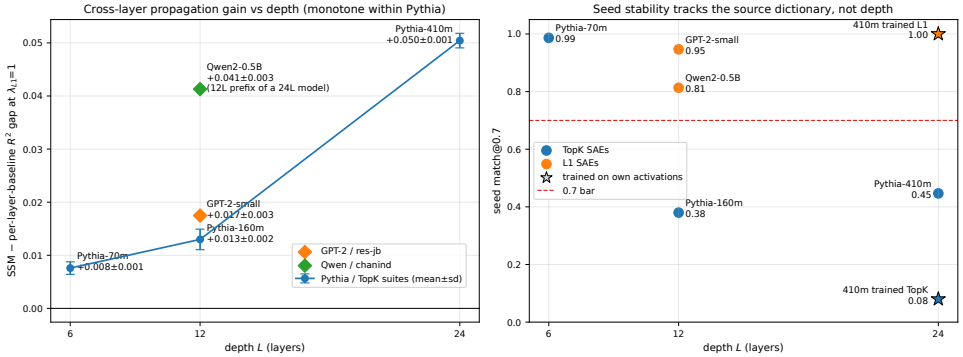


Figure 10: Cross-layer reconstruction gain (knee gap, SSM − matched per-layer baseline) versus depth: monotone within Pythia (3 seeds, mean±sd), replicated in sign at one point each in GPT-2 and Qwen2. Five caveats in Section 6.2 — mixed dictionary types; in-family depth confounded with width, parameter count and training-set size; a differently-constructed dataset at the 6-layer point; and a Qwen point that is a 12L SAE-coverage prefix of a 24L model (full-coverage value +0.072). The right panel plots decoder-direction seed stability (match@0.7) across the zoo, with the controlled within-model contrast starred: on the same Pythia-410m-deduped activations a trained L1 dictionary is perfectly stable (1.00) where trained TopK is not (0.08) — the dictionary, not the model or the depth, is what moves (Section 6.5).

points interpretability cares about. On GPT-2-small the two statistics agree instead (AR learned share 47%, reconstruction learned share 56%), so the dissociation is a one-family observation and we do not claim it more broadly. What carries across both families is the reconstruction result itself: roughly half the gain is learned.

6.3 Baselines that were not built for Δ, and it shows

The pre-registered indirect baselines both score *below zero* on update reconstruction: differencing the pretrained per-layer SAE's state reconstructions ($\hat{\Delta}_\ell = \hat{v}_{\ell+1} - \hat{v}_\ell$) gives $R^2 = \mathbf{-0.223}$ even though the same SAE reconstructs the *states* at +0.60 — the two reconstruction-error terms $\epsilon_{\ell+1} - \epsilon_\ell$ swamp the small update signal, the bias the protocol predicted, dramatic at 24 layers — and the Tuned-Lens-implied update gives $\mathbf{-0.309}$ (pooled; the layer-averaged version of the same baseline is the −0.691 of Section 3.2). The two Δ-targeted methods (SSM +0.867, matched per-layer dictionary +0.816, both at the sweep's densest point λ_{L1}=0.3) against the two indirect ones is the clearest statement of why the update needs to be modelled directly.

6.4 Against the unconstrained $O(L^2)$ parameterisation (CLT)

The primary conceptual comparison fits a *residual-write CLT* — a free low-rank decoder $W^{\ell' \to \ell}$ at every (source, target) pair, $\hat{\Delta}_\ell = b_\ell + \sum_{\ell' \le \ell} W^{\ell' \to \ell} u_{\ell'}$ with the same unpenalised per-layer bias the SSM gets, the unconstrained family our SSM restricts — on the same frozen sources at the same per-decoder rank, under the frozen-source protocol. Scope qualifiers stated once and inherited by every sentence below: residual-write variant (not MLP-write), frozen shared sources (CLT does not get its own dictionary), per-pair rank capped at 32–64 (rank 128 exceeds the study GPU; Section E), N = 1,244 training examples (1,057 where a validation carve-out is used), and no natively-trained CLT (limitation). The honest summary is three axes (Figure 11):

1. **Matched per-decoder rank: a small, real CLT edge.** At rank 32 the best-of-grid gap is SSM−CLT = −0.0142 [−0.0185, −0.0099]; re-selecting λ_{reg} on a train-internal validation split (never test) reproduces it at −0.0090 [−0.0136, −0.0045]. The edge exists at every depth tested — at the shared grid λ_{reg}=1 used for all models: 6L −0.009; 12L −0.050 GPT-2, −0.023 Pythia-160m; 24L −0.011; Qwen2's coarse-grid point, −0.021, agrees in sign but carries no CI — and is *not* monotone in depth; it tracks model and dictionary.

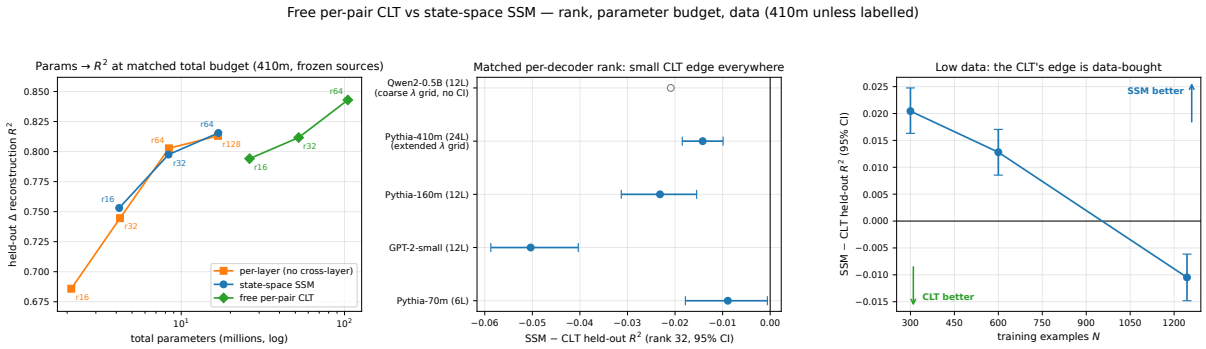


Figure 11: SSM vs residual-write CLT on frozen sources, three axes: (A) parameters → held-out R^2 at matched total budget (all decoder families near one frontier; per-layer rank-64/128 points included per Section 6.1); (B) per-model rank-matched gaps with paired CIs — the primary-model point is axis 1's extended-grid −0.0142, the other models sit on the shared λ_{reg}=1 grid, and Qwen2's coarse-grid point carries no CI (Section 6.4); (C) sample-efficiency sign flip (crossover $\approx 10^3$ examples).

(An earlier apparent tie was an artefact of a too-short regularisation grid; we report the corrected numbers. Widening the grid moves the primary model's CLT from 0.804 to 0.812, which is why the depth list above and axis 1's headline differ.)

2. **Matched total parameters: the SSM Pareto-dominates up to 52M.** Comparing each family at its best grid point, SSM rank-64 (16.9M parameters, R^2 0.8155) beats CLT rank-16 (26.2M, 0.794) *and* CLT rank-32 (52.4M, 0.8117); the CLT needs 105M (rank 64, 0.843) to lead outright. The rank-32 margin is thin (+0.004) and, being a cross-run comparison, carries no paired CI — what the axis supports is that the SSM matches or beats the CLT at ∼3× fewer parameters, not that it beats it by a lot. Per-parameter, the first-order state is the better spend in this regime.

3. **Sample efficiency: the CLT edge is data-bought.** Subsampling the train set flips the sign: at N=300 the SSM leads by +0.0204 [+0.0163, +0.0248], at N=600 by +0.0128, at N=1,244 the CLT leads by 0.0105 (the N-series holds one shared regularisation grid fixed at every N; the extended-grid and validation-selected full-N gaps are axis 1's −0.0142/−0.0090) — crossover at roughly 10^3 examples for 52M free parameters versus 8.4M shared ones.

Both cross-layer methods beat the no-propagation per-layer decoder at every depth tested (SSM−per-layer +0.022 to +0.052; CLT−per-layer +0.045 to +0.072; on the primary model SSM−per-layer +0.0518 [+0.0485, +0.0552] under validation-based selection). Read together with H2's failure, the conclusion is symmetric and, we think, useful in both directions: the free per-pair decoders are *validated* — they buy a small, significant, data-dependent margin, so nothing forces a shared kernel — and the first-order state-space compression captures most of the cross-layer reconstruction value at ∼16% of the parameters, dominating at matched budget and in low data.

6.5 Stability: the dictionary, not the model, controls identifiability

Gate 3 recorded one failed clause: top-100 decoder-direction stability across re-initialisations was 0.45 (match@0.7), below the pre-registered 0.7, while reconstruction was seed-identical — an equivalent-minima degeneracy, unmoved by 100× stronger operator regularisation and by shrinking d_z. Three findings close the story (Figure 12):

1. **It is the TopK sparsity mechanism.** Training matched L1 and TopK dictionaries on the *same* cached activations — identical architecture, matched L_0, differing only in the sparsity operator — and refitting: decoder match@0.7 — the fraction of the top-100 features whose

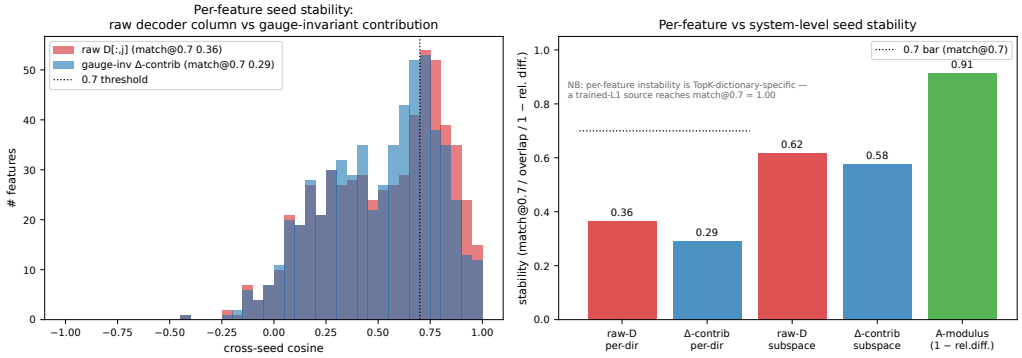


Figure 12: Per-feature vs system-level seed stability on the TopK (MLSAE) source. Left: cross-seed cosine of raw decoder columns (match@0.7 = 0.36) and of gauge-invariant per-feature Δ-contributions (0.29) — gauge-invariance does not rescue single-feature identifiability. Right: the system-level statistics the claims use are seed-stable (decoder subspaces; the spectrum bar plots the A_ℓ modulus agreement as $1-$median relative difference = 0.91 — the honest statistic of Section 6.8, not the vacuous sorted-spectrum cosine — and is a property of the fitted gauge, not a gauge invariant). The in-panel note carries the resolution: the per-feature instability is TopK-specific — a trained-L1 source on the same activations reaches match@0.7 = 1.00.

decoder direction has cross-seed cosine ≥ 0.7 — is **1.00 (L1)** versus **0.08 (trained TopK)** versus 0.45 (pretrained MLSAE TopK), at the same reconstruction win (+0.041–+0.047). Not a width effect: the two trained dictionaries have identical nominal width (8,192) and it is the *stable* one whose alive set is larger (1,178 L1 vs 660 TopK), so stability is not bought by narrowness. Cross-model, L1 sources are stable on two further families (GPT-2-small 0.95, Qwen2 0.81) while TopK is unstable on two Pythia models (0.38, 0.45), with N ruled out by subsampling; the zoo's one *stable* TopK point is the six-layer debug lane's EleutherAI k=16 suite (0.99, Section D), so the failure is not TopK-universal — but every TopK source we tested at $\geq$12 layers was unstable and every L1 source stable. Retraining the L1 dictionary on the train split only (no test leakage) changes nothing (match 1.00; gaps +0.037–+0.070).

2. **It is not a gauge artefact.** Gauge-invariant per-feature impulse responses (feedthrough plus downstream state path) are as unstable (0.29) as raw decoder columns (0.36 over the gauge analysis's 200-feature candidate set; the Gate-3 top-100 statistic quoted above is 0.45), so the instability is genuine single-feature non-identifiability over redundant TopK atoms — expected for a distributed code where the fit only constrains realised combinations.

3. **Everything the claims use is stable.** Reconstruction (0.880 every seed), the state-operator spectrum (median relative difference of the sorted moduli 0.086 across seeds — *not* the sorted-modulus cosine of 0.999, which random spectra also achieve; Section 6.8), the gauge-invariant state-path decay (gap-2 ratio 0.42–0.49 across seeds and dictionaries), the top-feature decoder *subspace* (∼0.6 overlap), and the feature *set* across data halves (Jaccard 0.89).

We flag the first finding as independently useful: pipelines that interpret *individual* SAE features while sourcing from TopK dictionaries are building on directions that need not survive a re-seed.

6.6 The causal existence proof: a clean, fully controlled null

The pre-registered critical test: find a seed-stable, interpretable feature, subtract its estimated contribution δ_ℓ from the residual stream during the real forward pass (re-solving the LASSO with

the frozen decoder, so reconstruction error stays untouched), and require — all four pre-registered — a ≥ 0.20 mean top-1 probability drop on the feature's top-100 examples, preserved controls, a small matched-random-feature effect, and a symmetric ≥ 1.0-nat injection effect on 70% of silent examples.

The test fails, cleanly. Across the six strongest candidates (from a pool of 53 seed-stable features with purity 1.00 and legible themes), the largest ablation effect is +0.024 — roughly an eighth of the 0.20 threshold — and the injection test returns −0.02 nats; controls and magnitude-matched random features are ≈ 0, so the splice mechanism is clean. This is the pre-registered honest landing place, and we report it as such.

The diagnostics then locate *why*, with an escalating control sequence (Figure 13): amplifying one feature's δ_ℓ ten-fold drives the answer probability most of the way to zero, saturating at full removal by 30× (direction correct, natural magnitude small); ablating each example's top-k *co-firing* features — pooled over eight seed-stable candidates, a superset of the six screened above — crosses threshold at k=100 (+0.237, CI $[0.213, 0.262]$, bootstrapped over the eight candidates, which are the independent unit here, not over test examples; the per-candidate means span 0.185–0.301, seven of the eight individually above threshold); a *count-matched* random group does not (+0.034) — but a *mass-matched* random group reproduces it (+0.244, i.e. slightly *exceeding* the real group), and, decisively, a **norm-matched random direction** — even an isotropic one in $\mathbb{R}^D$ — reproduces it too (+0.224 isotropic, +0.204 contribution-shaped, crossing at the same k). Against the isotropic direction, feature identity retains a significant but ∼5% sliver (+0.012 $[+0.004, +0.021]$); against the contribution-shaped direction +0.033 $[+0.020, +0.047]$; against the mass-matched *feature* group, −0.008. We quote the isotropic contrast because it is the control that isolates direction alone, and note that no control leaves feature identity with more than a seventh of the effect. Off-target specificity of the group ablation is weak and class-dependent: at the threshold-crossing group size it lowers the example's *own* answer's log-probability by 2.89 nats (pooled over candidates) against 2.45 nats for a basket of *other* classes' answer tokens (geography features answer-specific, organization features not). The group effect is therefore *magnitude-driven*: it demonstrates the query position's generic sensitivity to perturbations of that size, not a causally privileged code.

Swapping to the clean L1 dictionary strengthens rather than rescues the null: the candidate pool becomes 199/200 seed-stable, the maximum single-feature effect rises nearly five-fold to +0.112 — a less redundant dictionary *concentrates* causal mass — and still no feature approaches 0.20. **Factual recall in this model is a distributed, mass-driven code with no single-feature bottleneck, independent of the dictionary used to read it.** No threshold was adjusted post hoc.

6.7 The state is functionally informative, not reconstruction-only

Does the 32-dimensional state z_ℓ carry task information, or is it only reconstruction machinery? We probe four arms at matched dimension d=32 with identical linear probes (hyperparameters selected on a train-internal validation carve-out): the state; PCA-32 of the current layer's code u_ℓ; PCA-32 of the concatenated prefix codes $u_{\leq \ell}$ (an information superset of what the state sees); and PCA-32 of the raw residual v_ℓ as a reference. Target: 205-way answer identity (majority 0.106; permuted-label null at majority).

At mid-depth the state wins against both code summaries on both dictionaries (Figure 14): state−code +0.060 $[+0.037, +0.084]$ (MLSAE) / +0.139 $[+0.105, +0.173]$ (L1); state−prefix +0.080/+0.173 at mid-depth and +0.110/+0.188 late; at ℓ=17 — the answer-promotion depth where k^* blips (Figure 7) and per-layer reconstruction dips (Figure 16, right) — the state jumps from 0.57 to 0.76, a full layer before the code and raw-residual arms make their own jump, and beats even the raw-residual PCA there (0.759 vs 0.666). Peak state decoding reaches 0.775/0.801 top-1 from 32 dimensions (majority class 0.106; with 205 classes over 311 test examples, top-1 is dominated by the frequent objects, so we read the *contrasts between matched arms*, not the

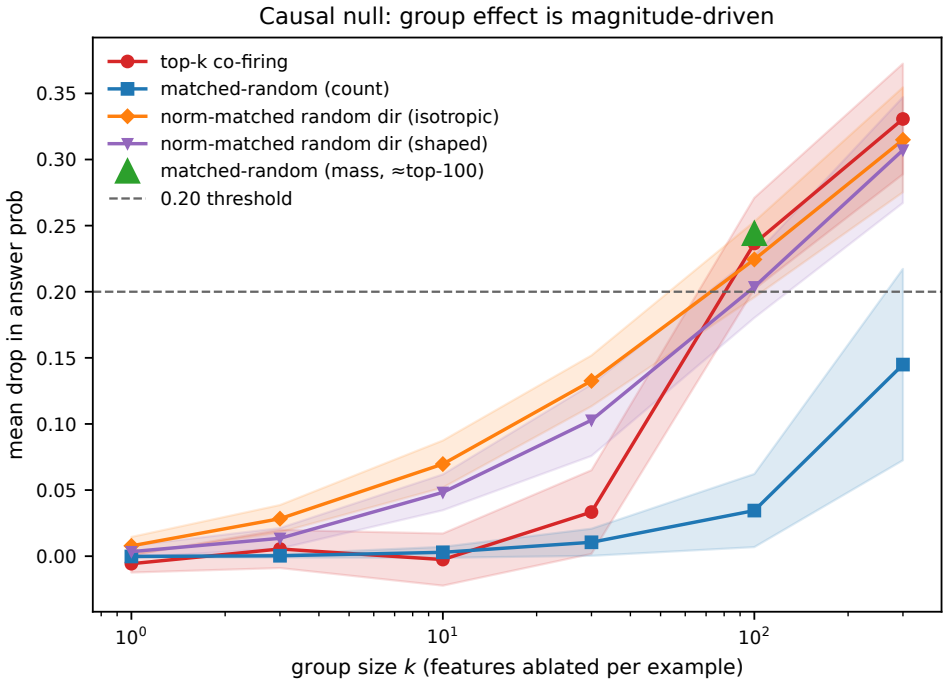


Figure 13: Group-ablation controls, pooled over 8 seed-stable candidates × 100 examples. Top-k co-firing ablation crosses the pre-registered 0.20 at k=100; count-matched random groups do not; mass-matched random groups and norm-matched random *directions* do — the effect tracks removed magnitude, not feature identity.

absolute level). Since the SSM is linear, the state and the prefix-PCA are two 32-dimensional *linear images of the same prefix codes*: the compression trained only to reconstruct Δ retains answer-identity information that the variance-optimal compression discards. Scope, stated plainly: the late band washes out against a per-layer code summary, the raw stream remains the best readout late (-0.04 to -0.07), relation class saturates for every arm, and model confidence favours the code arms — the state's advantage is specific to *what* is being recalled, at the depths where recall is assembled.

6.8 The learned state path attenuates fast

The state operators are the decomposition's structural readout. Getting the epistemics right here requires a gauge argument we initially got wrong, so we state it first.

What is and is not a property of the fit. Equation (1) has *free per-layer* A_ℓ, B_ℓ, C_ℓ. Its symmetry group is therefore a *per-layer* $T_\ell \in GL(d_z)$,

$$A_\ell \mapsto T_{\ell+1} A_\ell T_\ell^{-1}, \quad B_\ell \mapsto T_{\ell+1} B_\ell, \quad C_\ell \mapsto C_\ell T_\ell^{-1}, \quad z_0 \mapsto T_0 z_0,$$

which leaves every $\hat{\Delta}_\ell$ unchanged. Only a *global* $T_\ell \equiv T$ acts on A_ℓ by similarity, so **the eigenvalues of A_ℓ are not gauge-invariant**: even the scalar subgroup $T_\ell = c_\ell I$ rescales $\rho(A_\ell)$ by $c_{\ell+1}/c_\ell$. Applying a random output-preserving scalar gauge to our fitted checkpoint (predictions unchanged to 2×10^{-7}) moves the median ρ from 0.43 to 0.26, the maximum from 0.88 to 4.91, and the count of non-contractive transitions from 0/23 to 5/23. The quantity that *is* invariant — the T_ℓ telescope out of it — is the realised state path $C_\ell\,(A_{\ell-1} \cdots A_{j+1})\,B_j$ from source layer j to update ℓ. Under the same gauge it is unchanged to 1×10^{-8}.

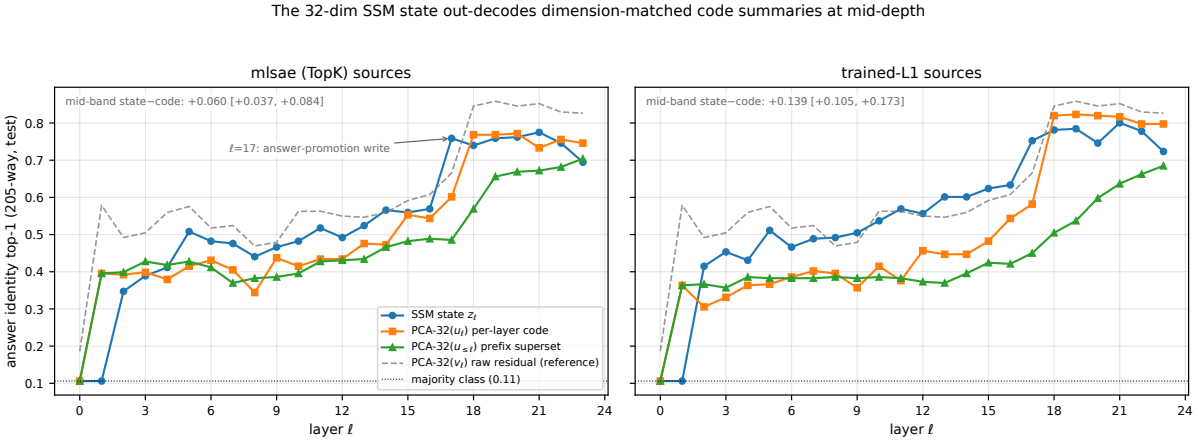


Figure 14: Answer-identity decoding (205-way top-1) per layer from four 32-dimensional summaries, both dictionaries. The Δ-reconstruction-trained state beats dimension-matched PCA of its own codes (and their prefix superset) at mid-depth on both dictionaries (the early-band edge is significant only under L1; under MLSAE it is null), and at ℓ=17 beats even raw-residual PCA; the raw stream remains best late.

The invariant result. The state-path kernel norm **halves per layer of depth gap** (Figure 15, right): normalised to gap 1, $1 \to \mathbf{0.485} \to 0.211 \to 0.089$ across gaps 1..4, an e-folding memory of $\tau = \mathbf{1.15}$ layers (τ is the least-squares log-linear fit over gaps 1–8; the per-gap ratio tightens from 0.49 at gap 2 to ∼0.40–0.44 at every longer gap, which is why the fitted τ sits slightly below the $1/\ln 2 \approx 1.44$ that exact halving would give). This is stable across three seeds and both dictionaries (gap-2 ratio 0.42–0.49, $\tau = 1.15$–1.21 layers). A source two layers back contributes half what one layer back does, three back a fifth — so first-order reads capture essentially all state value, and this is a *mechanism consistent with* $k^* \approx 1$. The epistemic caveat still holds: it is estimated on the same data on which k^* was measured, so it is an explanation within the model class, not an independent confirmation; the model-free evidence for $k^* \approx 1$ remains the per-depth AR sweeps of Section 4.2.

The gauge-dependent result, reported as such. In the minimum-Frobenius representative that the λ_{reg} penalty selects, the operator is contractive at every layer (Figure 15, left; spectral radius median 0.43, maximum 0.88, 0/23 transitions ≥ 1; the fit uses $A_0, \ldots, A_{22}$, since the transition out of the last state would feed no update). That is a reproducible property of *the regularised fit* — it recurs across seeds and dictionaries — and we report it because it is what one would look at, but no claim in this paper rests on it. We also retire one statistic we previously quoted: the cross-seed cosine of the sorted modulus spectrum is 0.999, but *two independent random spectra score* 0.999 *on the same statistic*, because both vectors are sorted and non-negative. The honest cross-seed number is the median relative difference of the sorted moduli, **0.086** (and 0.111 between dictionaries); the per-layer ρ *profile* correlates only 0.54 across dictionaries. So the modulus distribution is reproducible to about 10% and the layer assignment much less so.

Depth-varying (H2 fails) + a state path that halves per layer + a ∼1-layer memory is a compact, mechanistic summary of how these updates travel: *information written into the stream persists via the stream itself; the incremental dynamics forget almost immediately.*

6.9 Where the reconstruction lives: sources at every layer

Keeping sources only up to layer $\ell_{\max}$, refitting, and letting the SSM propagate them forward: held-out R^2 climbs monotonically over the sampled cutoffs $\ell_{\max} \in \{0, 3, 7, 11, 15, 19, 23\}$, from −0.44 (embedding-only) through +0.20 ($\ell_{\max}$=11) and +0.58 ($\ell_{\max}$=19) to +0.77 (full), still

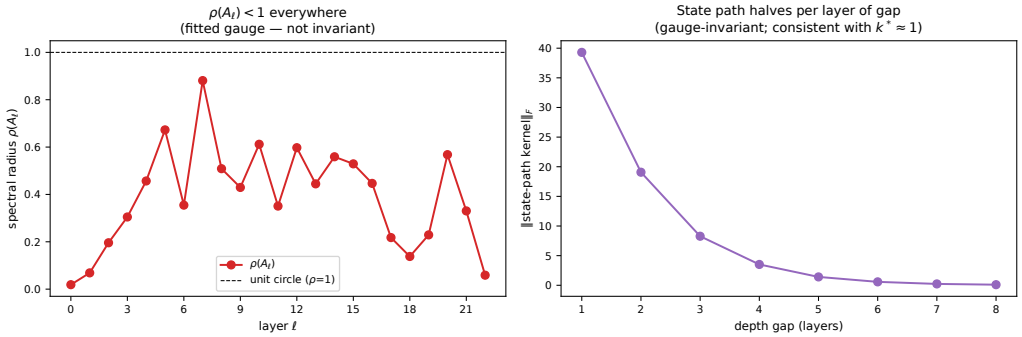


Figure 15: Left, *gauge-dependent*: in the fitted minimum-Frobenius gauge $\rho(A_\ell) < 1$ at every depth (median 0.43) — but a per-layer gauge change that leaves every prediction fixed sends 5/23 transitions above 1, so this panel describes the regularised fit, not the model. Right, *gauge-invariant*: the realised state-path kernel $\|C_\ell(\prod A)B_j\|_F$ halves per layer of depth gap ($\tau \approx 1.15$ layers), the operator-level mechanism consistent with $k^* \approx 1$ (Section 6.8).

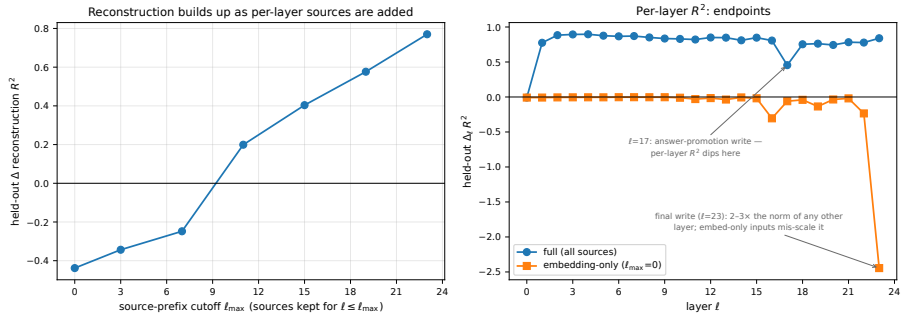


Figure 16: *Left*: source-prefix sweep — held-out R^2 as sources are admitted up to $\ell_{\max}$ (SSM free to propagate). Monotone over the seven sampled cutoffs, turning positive around mid-depth and still climbing at the top of the stack: late-layer sources are load-bearing and non-propagable. *Right*: per-layer R^2 at the sweep's two endpoints. The annotated interior dip at ℓ=17 (full-source arm: 0.46 against 0.81/0.75 at the adjacent layers) is the answer-promotion write — the same depth where k^* blips (Figure 7) and the state probe jumps (Figure 14). The dip is a feature of the per-layer panel, not of the prefix curve.

gaining +0.19 over the last four layers (Figure 16, left). The embedding-only floor is degenerate by construction — v_0 has no SAE, and this sweep runs under the frozen-source protocol, which does not re-infer codes, so here $u_0 \equiv 0$ and the state trajectory is example-independent (under the LASSO protocol u_0 is re-inferred like every other layer's code, Section 5.2) — and its −0.44 is almost entirely one layer: the per-layer R^2 stays small at 23 of 24 depths (within ±0.05 at 19 of them, never below −0.31) and falls to −2.4 at the final write, whose norm is 2–3× any other's. Late-layer sources carry reconstruction value that propagation *cannot* substitute — the features are genuinely new at depth, matching the staged-recall picture of Geva et al. (2023) — which justifies the design's per-layer sources and completes a two-sided ablation ladder with the state ablation: removing propagation costs the headline-sized margin (0.77 → 0.73 under this protocol — the state ablation of Section 6.1 seen from the other side); removing per-layer sources is catastrophic.

7 Limitations

Limitation 1 — feature novelty is structurally out of reach. The pre-registered novelty test (find a feature no baseline finds) reduces to a structural fact here: the SSM sources from the *same* SAE atoms as the per-layer baselines, so no feature can be novel by activation identity. The contribution is the cross-layer propagation structure over those atoms; no novelty fraction is

computable or claimed.

Limitation 2 — cross-family coverage stops below 1B parameters. Every experiment ran on one 3 GB GPU (GTX 1050; Section E), which admits Pythia to 410m, GPT-2 to large, and Qwen2-0.5B, but not Gemma-2-2B or Llama-3.1-8B. The H1/H2 replication evidence is therefore three families at 12–36 layers (the six-layer Pythia-70m lane is a depth point for the Phase-3 fit only; its differently constructed dataset keeps it outside the H1/H2 claim, Section 3.1), and the two questions this leaves open are exactly the ones a reader should ask. First, *does any of this hold at 7B?* We do not know, and we state the expectation as a falsifiable prediction: the sign pattern (H1 above null, H2 cross-band negative, $k^* \approx 1$) should persist at scale, since it held without exception across the six models spanning 124M–774M parameters and depths 12–36 in three families; we expect the fast-attenuating state path to persist as well, though it is measured on the primary model only (Section 6.8); the magnitudes (in particular the depth-graded reconstruction gain) we make no prediction for. Second, *within a family, depth is confounded with width, parameter count and training-set size* on the depth-scaling series (Section 6.2). One contrast partially deconfounds it: the Qwen2 prefix comparison holds the model, the activations, the dataset and the dictionary type fixed and varies only how many layers the fit covers (24 vs 12), and the gain rises with covered depth (+0.072 vs +0.043) — depth doing work at constant width. It is one contrast, on one model, without an interval. A clean in-family resolution (Pythia 1b/1.4b/2.8b, one dictionary type, matched N, three seeds) is specified and costed as future work, along with the two missing families.

Limitation 3 — three relation classes. The primary model confidently knows too few CounterFact facts for the protocol's original ≥1,000-per-class target, so the study is scoped to geography, language, and organization. No load-bearing claim depends on class breadth — the reconstruction and characterisation results are cross-model-replicated and the causal result is a controlled null — but cross-relation generality beyond these classes is untested, and the observed dissociation (organization features reconstruct class-specifically yet ablate non-specifically; geography features are specific on both) is reported on three classes without a breadth claim.

Scope of the characterisation. All kernels are fit at the *query position* of factual-recall prompts; cross-position transfer is negative (Section 4.3), so nothing here describes token positions in general. The generic-text replication uses short windows (median 9 content tokens) length-matched to the prompts. The attention/MLP split of Section 3.5 is exact only for parallel-block architectures like Pythia. Correctness-based probe targets are degenerate by construction (the dataset is filtered to model-correct facts), which is why the probe targets answer identity.

Methodological. The held-out protocol is reconstruction-at-matched-sparsity (codes re-inferred on test Δ with frozen operators), standard for dictionary learning but not out-of-sample forecasting; all arms are treated identically and Section B tabulates the operating points. Some hyperparameter selection touched the test split: AR best-(k, rank) is a max over symmetric real-vs-null evaluations (the contrast is the claim), and the CLT family's λ_{reg} was re-selected on a validation carve-out with the conclusion unchanged (Section 6.4); the Granger F approximation is empirically calibrated rather than assumed (Section C). Worse, that AR max ranges over fits scored on *different depth sets* ($\ell \geq k$), so a larger k can win by landing on an easier subset; Section 3.3 states the confound and Section C re-derives every trained-vs-untrained and factual-vs-generic contrast at a common k=1. The pre-registered headline number is left as first reported. Finally, the fitted state operator's spectral radius is *not* a gauge-invariant quantity in a depth-indexed SSM (Section 6.8); we report it as a property of the minimum-Frobenius representative the regulariser selects, and rest the mechanism claim on the invariant state-path

decay instead. An earlier version of this work asserted the invariance, and the correction is one of the reasons the operator section is phrased the way it is. The CLT comparison is the residual-write variant on frozen shared sources with per-pair rank ≤ 64; a natively-trained, jointly-sourced CLT could behave differently and remains the right off-box follow-up. Untrained-network controls use a single random initialisation each. Pythia's AR conclusion is far from any decision boundary (the untrained network *exceeds* the trained one), but the quantified shares — GPT-2-small's 47% learned AR structure and the ≈half-learned reconstruction gains — are single-initialisation point estimates with no seed-level spread attached, so we hold them at "about half" and claim nothing finer. Block-Toeplitz structure was always an approximation to test, not a derivation from the nonlinear architecture; H2's failure retires it.

Interpretive. H1's headline numbers are partly tautological and substantially architectural (Section 3.4); we consider the quantified version of that sentence a contribution, but it means raw AR-on-Δ levels from this or any similar pipeline should never be quoted as learned depth-computation. The causal null is a null for *this* feature-ablation methodology at *these* thresholds on *this* circuit family; it constrains single-feature stories of recall, not all causal uses of sparse decompositions.

8 Discussion

What kind of thing is depth structure? The composite picture is specific enough to be useful. Residual updates are predictable across depth, but the predictability decomposes into (i) a running-sum persistence term that is definitionally there, (ii) a genuine increment-level structure that a randomly initialised network already exhibits — substantially so on Pythia-410m-deduped, where the untrained network in fact exceeds the trained one; about half of it on GPT-2-small — and (iii) a learned component that only becomes visible in the *generative* fit — roughly half the SSM's reconstruction advantage, majority at sparse operating points. Meanwhile the dynamics are pervasively non-stationary in depth, first-order, and — per the fitted model's gauge-invariant state path — attenuating by half per layer of gap, a one-layer memory. A transformer, on this evidence, does not implement a depth-shared propagation rule; it implements a different small-memory rule at every layer, over features that persist via the additive stream itself rather than via the incremental dynamics. The Granger reach of 20 layers coexisting with $k^* \approx 1$ stops being a paradox once split by sublayer: a small, deep-reaching attention channel superposed on an energy-dominant, memoryless MLP channel. (That split localises the reach; it does not by itself certify the AR structure as persistence-free, since both sublayer writes read the same drifting state — Section 3.5.)

Consequences for cross-layer interpretability methods. H2's failure is an independent validation of CLT's design decision to give every (source, target) pair free parameters (Ameisen et al., 2025) — our attempt to compress that family along depth was the pre-registered falsifier, and it failed in CLT's favour. The quantitative comparison then adds the budget-conditional half of the story: the free decoders buy a small, real, data-bought margin, while a first-order state captures most of the cross-layer value at a sixth of the parameters and generalises better below 10^3 examples. For practitioners the operational reading is: if you need cross-layer reconstruction under a parameter or data budget, a depth-indexed state is the better spend; if you need maximum fidelity and have data, free per-pair decoders are not redundant. Our TopK finding adds an orthogonal warning: single-feature decoder directions from TopK dictionaries were seed-unidentifiable in every configuration we tested beyond the six-layer debug lane, while matched L1 dictionaries were stable everywhere (perfectly so on the primary model) — any pipeline that names, steers, or ablates *individual* TopK features should check its features survive a re-seed.

Consequences for factual-recall stories. The causal null is not a failure to find structure; it is a measurement of what kind of structure recall has. Single features move the answer by at most an eighth of the pre-registered threshold under the TopK dictionary, and by about half of it under the cleaner L1 dictionary; the effect that does clear threshold is reproduced by a norm-matched random direction; and the cleaner dictionary concentrates but does not localise the effect. This is a controlled, dictionary-independent statement that recall at the query position is a distributed, magnitude-dominated code — consistent with the drift in the editing literature away from single-component accounts (Hase et al., 2023), and a calibration point for ablation-based evaluation generally: group-ablation effects need magnitude-matched *direction* controls before they can be read as evidence about privileged feature sets.

Consequences for probing. The state-probe result cuts the other way: trained purely to reconstruct updates, the 32-dimensional state retains answer-identity information that variance-optimal compressions of the same codes discard. Reconstruction objectives, at least here, learn compressions aligned with what the circuit is computing — a modest but genuinely positive datum for the program of extracting functional summaries from decompositions, and one that pairs naturally with the fast-attenuating state path: a short-memory, information-dense running summary is exactly what a fast-forgetting state can be.

The EMG intuition, in retrospect. The convolutive-source-separation analogy (Holobar and Zazula, 2007) supplied H2, and H2 is false: transformer depth is not a homogeneous propagation medium. What survives from the analogy is the discipline it imported — model the increments, put sources at every position, and treat the mixing structure as an estimable, testable object. The estimated object just turned out to be depth-indexed, fast-attenuating, and circuit-specific rather than shared.

Future work. Three items are specified and costed in the project record: deconfounding the depth series from width (Pythia-1b/1.4b/2.8b with one dictionary type, three seeds); the two missing families (Gemma-2's sandwich normalisation, Llama-3.1-8B) which also enable the pre-registered $\geq$1,000-per-class relation test; and a natively-trained CLT. Two more suggest themselves from the results: whether the state-path decay rate and the mid-depth functional-state window are universal across scale, and whether the sublayer AR structure survives the T2 surrogate and the untrained control that Section 3.5 did not run.

9 Conclusion

We pre-registered a hypothesis about how residual-stream updates travel through transformer depth, and the data answered with unusual precision: structured, yes — but the structure is substantially architectural where it is predictive (on the primary model; about half learned on GPT-2-small), pervasively depth-indexed where we hoped it was shared, first-order and fast-forgetting where it is generative, distributed and magnitude-driven where we tested it causally, and genuinely learned exactly where the decomposition adds value over per-layer baselines. Both pre-registered hypotheses did real work: H1 by surviving controls designed to kill it, H2 by failing and thereby validating the unconstrained parameterisation the field already uses. Every gate decision, threshold, and fallback was written down in advance; every headline reconstruction gap carries a paired bootstrap confidence interval (the depth series carries seed error bars instead, and Section 6 names the handful of supporting numbers that carry no interval at all); the two places where our own analysis misled us — a spectral radius we wrongly called gauge-invariant, and an autoregressive sweep whose depth set moved with the lag order — are corrected in the open, with the statistics that survive them; and the full pipeline — from cached activation tables

to every figure in this paper — regenerates with one script. We offer the results, the nulls, and the protocol in equal measure.

Acknowledgements

Drafting, code review, and experiment execution were assisted by AI coding tools; all analyses, decisions, and text were reviewed by the author.

References

Emmanuel Ameisen, Jack Lindsey, Adam Pearce, Wes Gurnee, Nicholas L. Turner, Brian Chen, Craig Citro, David Abrahams, Shan Carter, Basil Hosmer, Jonathan Marcus, Michael Sklar, Adly Templeton, Trenton Bricken, Callum McDougall, Hoagy Cunningham, Thomas Henighan, Adam Jermyn, Andy Jones, Andrew Persic, Zhenyi Qi, T. Ben Thompson, Sam Zimmerman, Kelley Rivoire, Thomas Conerly, Chris Olah, and Joshua Batson. Circuit tracing: Revealing computational graphs in language models. *Transformer Circuits Thread*, 2025. URL https://transformer-circuits.pub/2025/attribution-graphs/methods.html.

Amir Beck and Marc Teboulle. A fast iterative shrinkage-thresholding algorithm for linear inverse problems. *SIAM Journal on Imaging Sciences*, 2(1):183–202, 2009.

Nora Belrose, Zach Furman, Logan Smith, Danny Halawi, Igor Ostrovsky, Lev McKinney, Stella Biderman, and Jacob Steinhardt. Eliciting latent predictions from transformers with the Tuned Lens. *arXiv preprint arXiv:2303.08112*, 2023.

Stella Biderman, Hailey Schoelkopf, Quentin Anthony, Herbie Bradley, Kyle O'Brien, Eric Hallahan, Mohammad Aflah Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar van der Wal. Pythia: A suite for analyzing large language models across training and scaling. In *Proceedings of the 40th International Conference on Machine Learning (ICML)*, 2023.

Joseph Bloom. Open source sparse autoencoders for all residual stream layers of GPT2-Small, 2024. AI Alignment Forum; weights at `jbloom/GPT2-Small-SAEs-Reformatted`.

Trenton Bricken, Adly Templeton, Joshua Batson, Brian Chen, Adam Jermyn, Tom Conerly, Nick Turner, Cem Anil, Carson Denison, Amanda Askell, Robert Lasenby, Yifan Wu, Shauna Kravec, Nicholas Schiefer, Tim Maxwell, Nicholas Joseph, Zac Hatfield-Dodds, Alex Tamkin, Karina Nguyen, Brayden McLean, Josiah E. Burke, Tristan Hume, Shan Carter, Tom Henighan, and Christopher Olah. Towards monosemanticity: Decomposing language models with dictionary learning. *Transformer Circuits Thread*, 2023. URL https://transformer-circuits.pub/2023/monosemantic-features/index.html.

David Chanin. Residual-stream sparse autoencoders for Qwen2-0.5B, 2024. Weights at `chanind/sae-qwen2-0.5b-res` (Hugging Face).

Hoagy Cunningham, Aidan Ewart, Logan Riggs, Robert Huben, and Lee Sharkey. Sparse autoencoders find highly interpretable features in language models. In *The Twelfth International Conference on Learning Representations (ICLR)*, 2024. arXiv:2309.08600.

Damai Dai, Li Dong, Yaru Hao, Zhifang Sui, Baobao Chang, and Furu Wei. Knowledge neurons in pretrained transformers. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2022.

Yanai Elazar, Nora Kassner, Shauli Ravfogel, Abhilasha Ravichander, Eduard Hovy, Hinrich Schütze, and Yoav Goldberg. Measuring and improving consistency in pretrained language models. *Transactions of the Association for Computational Linguistics*, 9:1012–1031, 2021.

EleutherAI. Per-layer sparse autoencoders for Pythia-70m-deduped, 2024. Weights at `EleutherAI/sae-pythia-70m-deduped-32k` (Hugging Face).

Nelson Elhage, Neel Nanda, Catherine Olsson, Tom Henighan, Nicholas Joseph, Ben Mann, Amanda Askell, Yuntao Bai, Anna Chen, Tom Conerly, Nova DasSarma, Dawn Drain, Deep Ganguli, Zac Hatfield-Dodds, Danny Hernandez, Andy Jones, Jackson Kernion, Liane Lovitt, Kamal Ndousse, Dario Amodei, Tom Brown, Jack Clark, Jared Kaplan, Sam McCandlish, and Chris Olah. A mathematical framework for transformer circuits. *Transformer Circuits Thread*, 2021. URL https://transformer-circuits.pub/2021/framework/index.html.

Nelson Elhage, Tristan Hume, Catherine Olsson, Nicholas Schiefer, Tom Henighan, Shauna Kravec, Zac Hatfield-Dodds, Robert Lasenby, Dawn Drain, Carol Chen, Roger Grosse, Sam McCandlish, Jared Kaplan, Dario Amodei, Martin Wattenberg, and Christopher Olah. Toy models of superposition. *Transformer Circuits Thread*, 2022. URL https://transformer-circuits.pub/2022/toy_model/index.html.

Lucy Farnik, Tim Lawson, Conor Houghton, and Laurence Aitchison. Jacobian sparse autoencoders: Sparsify computations, not just activations. In *Proceedings of the 42nd International Conference on Machine Learning (ICML)*, 2025. arXiv:2502.18147.

David Freedman and David Lane. A nonstochastic interpretation of reported significance levels. *Journal of Business & Economic Statistics*, 1(4):292–298, 1983.

Leo Gao, Stella Biderman, Sid Black, Laurence Golding, Travis Hoppe, Charles Foster, Jason Phang, Horace He, Anish Thite, Noa Nabeshima, Shawn Presser, and Connor Leahy. The Pile: An 800GB dataset of diverse text for language modeling. *arXiv preprint arXiv:2101.00027*, 2020.

Leo Gao, Tom Dupré de la Tour, Henk Tillman, Gabriel Goh, Rajan Troll, Alec Radford, Ilya Sutskever, Jan Leike, and Jeffrey Wu. Scaling and evaluating sparse autoencoders. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025. arXiv:2406.04093.

Mor Geva, Jasmijn Bastings, Katja Filippova, and Amir Globerson. Dissecting recall of factual associations in auto-regressive language models. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.

Peter Hase, Mohit Bansal, Been Kim, and Asma Ghandeharioun. Does localization inform editing? Surprising differences in causality-based localization vs. knowledge editing in language models. In *Advances in Neural Information Processing Systems 36 (NeurIPS)*, 2023.

Aleš Holobar and Damjan Zazula. Multichannel blind source separation using convolution kernel compensation. *IEEE Transactions on Signal Processing*, 55(9):4487–4496, 2007.

Diederik P. Kingma and Jimmy Ba. Adam: A method for stochastic optimization. In *The Third International Conference on Learning Representations (ICLR)*, 2015.

Tim Lawson, Lucy Farnik, Conor Houghton, and Laurence Aitchison. Residual stream analysis with multi-layer SAEs. In *The Thirteenth International Conference on Learning Representations (ICLR)*, 2025. arXiv:2409.04185.

Jack Lindsey, Adly Templeton, Jonathan Marcus, Thomas Conerly, Joshua Batson, and Christopher Olah. Sparse crosscoders for cross-layer features and model diffing. *Transformer Circuits Thread*, 2024. URL https://transformer-circuits.pub/2024/crosscoders/index.html.

Alireza Makhzani and Brendan Frey. k-sparse autoencoders. *arXiv preprint arXiv:1312.5663*, 2013.

Kevin Meng, David Bau, Alex Andonian, and Yonatan Belinkov. Locating and editing factual associations in GPT. In *Advances in Neural Information Processing Systems 35 (NeurIPS)*, 2022.

Neel Nanda and Joseph Bloom. TransformerLens, 2022. URL https://github.com/TransformerLensOrg/TransformerLens.

Adam Paszke, Sam Gross, Francisco Massa, Adam Lerer, James Bradbury, Gregory Chanan, Trevor Killeen, Zeming Lin, Natalia Gimelshein, Luca Antiga, Alban Desmaison, Andreas Köpf, Edward Yang, Zachary DeVito, Martin Raison, Alykhan Tejani, Sasank Chilamkurthy, Benoit Steiner, Lu Fang, Junjie Bai, and Soumith Chintala. PyTorch: An imperative style, high-performance deep learning library. In *Advances in Neural Information Processing Systems 32 (NeurIPS)*, 2019.

Fabian Pedregosa, Gaël Varoquaux, Alexandre Gramfort, Vincent Michel, Bertrand Thirion, Olivier Grisel, Mathieu Blondel, Peter Prettenhofer, Ron Weiss, Vincent Dubourg, Jake Vanderplas, Alexandre Passos, David Cournapeau, Matthieu Brucher, Matthieu Perrot, and Édouard Duchesnay. Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12:2825–2830, 2011.

Alec Radford, Jeffrey Wu, Rewon Child, David Luan, Dario Amodei, and Ilya Sutskever. Language models are unsupervised multitask learners. Technical report, OpenAI, 2019.

Skipper Seabold and Josef Perktold. statsmodels: Econometric and statistical modeling with Python. In *Proceedings of the 9th Python in Science Conference*, 2010.

Wei Shi, Sihang Li, Tao Liang, Mingyang Wan, Guojun Ma, Xiang Wang, and Xiangnan He. Route sparse autoencoder to interpret large language models. In *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025. arXiv:2503.08200.

Adly Templeton, Tom Conerly, Jonathan Marcus, Jack Lindsey, Trenton Bricken, Brian Chen, Adam Pearce, Craig Citro, Emmanuel Ameisen, Andy Jones, Hoagy Cunningham, Nicholas L. Turner, Callum McDougall, Monte MacDiarmid, C. Daniel Freeman, Theodore R. Sumers, Edward Rees, Joshua Batson, Adam Jermyn, Shan Carter, Chris Olah, and Tom Henighan. Scaling monosemanticity: Extracting interpretable features from Claude 3 Sonnet. *Transformer Circuits Thread*, 2024. URL https://transformer-circuits.pub/2024/scaling-monosemanticity/index.html.

An Yang, Baosong Yang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Zhou, Chengpeng Li, Chengyuan Li, Dayiheng Liu, Fei Huang, et al. Qwen2 technical report. *arXiv preprint arXiv:2407.10671*, 2024.

A Pre-registration and deviations

The protocol was written as four frozen phase documents (diagnostic, characterisation, decomposition, evaluation) plus a research statement, before any data were collected; a separate progress tracker was the only mutable record, so no pre-registered bar could be edited after the fact. The registration is internal rather than platform-hosted, but it is auditable: the four phase documents are plain files in the released repository, present in its initial commit and thereafter edited exactly once, to remove a process line (no quantitative content), so their state at every analysis step is verifiable from the public git history. The gates:

- **Gate 1** (H1 supported above null?) — *PASS*, all six pre-registered questions answered yes; routed to Phase 2.

- **Gate 2** (H2 stationarity?) — *FAIL (pervasive drift)*; routed, per the pre-registered decision matrix, to the state-space branch with free per-layer operators.
- **Gate 3** (reconstruction competitive at matched sparsity, and seed-stable features?) — *GO*: reconstruction clause met with margin; the seed-stability clause initially failed ($0.45 < 0.7$) and was recorded as a limitation per the plan's "reconstruction fine / seed unstable" branch, later closed by the dictionary-mechanism finding of Section 6.5.
- **Causal existence proof** — *NULL* at all four pre-registered thresholds; reported at the plan's designated honest landing place, with no threshold adjusted.

Deviations from the frozen protocol, exhaustively:

1. *Dataset thresholds* (plan-sanctioned tuning direction): CounterFact confidence threshold $0.3 \to 0.1$ and per-class floor $500 \to 300$, because the primary model knows only ∼2k facts; yield 1,555 examples / 3 classes instead of the aspirational 5–10k / ≥4 classes (whose ≥1,000-per-class *target* is the figure Section 7 cites; the 500 was the per-class inclusion *floor*). Relation-class scope became Limitation 3 by an explicit recorded decision (option (b) of the scale-up fork) after the cross-model replications made class breadth non-load-bearing.
2. *Primary checkpoint*: Pythia-410m → Pythia-410m-*deduped* (the public Tuned Lens required for the pre-registered gate baseline exists only for deduped).
3. *Stationarity-test order cap* (in-plan): $k = 5$ from the plan's early-band feasibility rule; working rank promoted $128 \to 256$ by the plan's own r-vs-$2r$ spot-check.
4. *Cross-model targets*: Gemma-2-2B / Llama-3.1-8B do not fit the 3 GB study GPU; replaced by the five replication models of Table 1 that do (with the primary: six models, three families; Limitation 2). The plan's Tuned-Lens gate was waived-and-reported on replication models without public lenses.
5. *Statistic repair* (post-hoc, both statistics archived): the depth-series gap statistic changed from max-over-λ matched-L_0 interpolation to the raw same-λ knee gap after per-seed forensics showed the former pairs against a saturated baseline tail; the knee statistic is the conservative one.
6. *Post-hoc analyses* beyond the protocol are labelled as such throughout (the T1–T5 pressure battery — T4 is its untrained-network AR control — and the untrained reconstruction analogue, the sublayer split, parameter-matched baselines, the state probe, and the review-driven CI/validation-selection reruns).
7. *Two self-corrections*, made after the draft was written and reported in the open because both had propagated into headline sentences. (a) The fitted state operator's spectral radius was described as gauge-invariant; in a depth-indexed SSM it is not, and the mechanism claim now rests on the invariant state-path decay (Section 6.8). (b) The pooled AR(k) sweep scores a depth set that shrinks with k, so "best over k" silently changed the evaluation set; the affected contrasts are re-derived at a common k in Table 3, and the pre-registered headline is reported as the four-depth number it is. Neither correction changes a gate outcome.

B Operating points and evaluation protocols

All held-out R^2 values are against the per-depth-mean baseline on the same 311-example test split. The apparent spread of "the SSM's R^2" across sections is protocol, not inconsistency:

Table 2: Every SSM (and matched-baseline) operating point quoted in this paper; "baseline" rows are the state-ablated twin (`use_state=False`). Gaps are differenced at full precision and then rounded, so a row pair is quoted to as many decimals as its gap needs.

Value	Protocol	Setting	Used for
0.880	train fit	λ_{L1}=1, r=d_z=32	convergence/sanity
0.867	LASSO held-out	λ_{L1}=0.3	Pareto top end (Figure 9)
0.857	LASSO held-out	λ_{L1}=1 (knee)	headline gap +0.049
0.698	LASSO held-out	λ_{L1}=30	Pareto bottom end (Figure 9)
0.849	LASSO held-out	rank-64 baseline, λ_{L1}=1	param-matched gap +0.008
0.816	LASSO held-out	baseline, λ_{L1}=0.3	baseline Pareto top end
0.808	LASSO held-out	baseline, λ_{L1}=1 (knee)	knee baseline of the +0.049 gap
0.653	LASSO held-out	baseline, λ_{L1}=30	baseline Pareto bottom end
0.79335	frozen-source	SSM rank 32, shared grid λ_{reg}=1	CLT depth list, N series
0.80383	frozen-source	CLT rank 32, shared grid	the −0.0105 gap
0.7975 / 0.8117	frozen-source	SSM / CLT rank 32, extended grid	the −0.0142 gap (axis 1)
0.7879 / 0.7969	frozen-source	SSM / CLT rank 32, validation-selected	the −0.0090 gap (N=1,057)
0.8155 / 0.843	frozen-source	SSM rank 64 / CLT rank 64	CLT budget axis (Section 6.4)
0.770	frozen-source	rank 32, $\ell_{\max}$=23	prefix sweep (Figure 16)

The LASSO protocol re-infers test codes against test Δ with frozen operators (reconstruction at matched sparsity; identical for both arms of every gap). The frozen-source protocol fits decoders on fixed cached codes and isolates decoder parameterisation; its levels sit uniformly lower. Gaps are only ever computed within a protocol.

C Statistical details

Granger-style nested-F. At each (depth ℓ, lag j) cell we compare nested AR models (lags $1..j-1$ vs $1..j$, per-depth intercepts included) on a top-32 PCA projection of Δ (full-D Wilks is singular at $D \gg N$), using Wilks' Λ with Rao's F approximation. Bonferroni corrects across the 112 feasible cells ($\alpha = 4.5 \times 10^{-4}$ per cell); the feasibility count follows the plan's rule ($\ell \geq j$), i.e. $\sum_j (L - j)$ over the lag grid. With the full grid $\{1, 2, 3, 5, 10, 15, 20\}$ this gives 112 cells at L=24 and 196 at L=36; the 12-layer replications (Table 1) cap the grid at $\{1, 2, 3, 5, 8, 10\}$ — lags 15 and 20 are infeasible at L=12, with lag 8 substituted to keep six grid points — giving 43 cells, with Bonferroni taken over the per-model feasible count.

Permutation calibration. Because the rogue high-variance dimensions make errors heavy-tailed, the F thresholds were checked with a Freedman–Lane permutation scheme (Freedman and Lane, 1983): fit the reduced model, permute its residuals, recompute F (999 permutations), at 12 spot cells spanning mid/late depth and lags 1–20. All 12 observed F statistics sit at the permutation floor ($p \leq 10^{-3}$), and the permutation null's tail quantiles match the parametric $F(\text{df}_1, \text{df}_2)$ within 1% at the 95th/99th percentiles — the Bonferroni cut behaves as advertised under exchangeability. The scrambled-null grids (0/112 significant, both stores) are the complementary full-grid empirical calibration.

Paired bootstrap CIs. Every headline gap is a paired bootstrap over the 311 test examples (10k resamples) on per-example sums of squares, computed from saved per-example arrays — pure post-processing, no refitting. Seed-level uncertainty, where reported (Figure 10), is mean±sd over three full refits.

Selection, and the AR depth-set confound. AR best-(k, rank) maximises held-out R^2 over 28 symmetric real-vs-null evaluations; the real−null contrast is the claim. But a pooled AR(k) fit scores only depths $\ell \geq k$, so the seven k values are evaluated on seven different depth

Table 3: The T1/T5 pressure battery re-run at a *common* lag order (k=1, every feasible depth), so trained/untrained and factual/generic are scored on the same targets. $R^2_{m=0}$ is the raw pooled AR(1) level; $R^2_{m=50}$ strips the top-50 persistent state directions (T1); the last two columns are T5's exact orthogonal split of $R^2_{m=0}$ at m=50 (tautological + genuine = raw). Scrambled nulls are $\leq 7 \times 10^{-3}$ in every row. Run `r13-common-depth-pressure`.

Store	depths	$R^2_{m=0}$	$R^2_{m=50}$ (T1)	T5 taut.	T5 genuine
Pythia-410m-deduped factual, trained	23	0.521	**0.436**	0.272	0.248
Pythia-410m-deduped factual, untrained	23	0.543	**0.503**	0.116	0.427
Pythia-410m-deduped generic, trained	23	0.882	**0.069**	0.879	0.003
Pythia-410m-deduped generic, untrained	23	0.063	**0.061**	0.007	0.056
GPT-2-small factual, trained	11	0.626	**0.439**	0.490	0.136
GPT-2-small factual, untrained	11	0.246	**0.233**	0.046	0.200

sets (23, 22, 21, 19, 14, 9 and 4 target depths at L=24). Selecting k by held-out R^2 therefore partly selects the easiest depth subset, and the pre-registered headline ($R^2 = 0.65$ at k=20) is an average over $\ell \in [20, 23]$, not over the stack. Two consequences we correct rather than hide:

- *Isolating the lag effect.* Re-running the k sweep with the depth set held fixed at $\ell \in [20, 23]$ separates lag order from depth subset, and agrees with the per-depth verdict of Section 4.2: order, not reach, is what reconstruction sees.
- *Like-for-like contrasts.* The T1 m sweep and the T5 split are re-run at fixed k=1 (all feasible depths) on every store the pressure battery uses. Table 3 is that table; all trained-vs-untrained and factual-vs-generic comparisons in Section 3.4 are taken from it. The pre-registered headline numbers are reported unchanged alongside. One conclusion changed: at a common k, GPT-2-small's tautology-free AR structure is about half learned, not ∼5%, so the "architectural" reading of T4 is specific to Pythia-410m-deduped.

Three readings follow. (i) On Pythia-410m-deduped the untrained network's persistence-stripped R^2 *exceeds* the trained network's (0.503 > 0.436), so T4's architectural verdict stands on the primary model. (ii) On GPT-2-small it does not (0.439 trained vs 0.233 untrained; learned share 47%). (iii) Generic text, compared at the same k, retains an order of magnitude less tautology-free structure than factual text (0.069 vs 0.436) — the raw generic level of 0.88 is 99.7% drift.

The CLT family's λ_{reg} was originally selected on test (house convention, interior peaks); Section 6.4 re-derives the one consequential CI with validation-based selection, where every arm picks an interior λ^*=3 and the conclusion is unchanged (test selection had mildly flattered the CLT, not the SSM). The four CLT gap numbers in Section 6.4 map to protocols as follows: −0.0105 is the shared grid (λ_{reg}=1, the grid every cross-model point uses); −0.0142 is the extended grid, best-of-grid per arm; −0.0090 is validation-selected on a 1,057-example train carve-out; −0.011 in the depth list is the shared-grid primary-model point, i.e. −0.0105 rounded.

Statistics defined once. *match@0.7* is the fraction of a run's top-n features (by activation mass; $n = 100$ at Gate 3, $n = 200$ in the gauge analysis) whose decoder direction has cosine ≥ 0.7 with its best match in a re-seeded fit. *Purity* is the largest relation class's share of a feature's top-100 activating examples. *Retention ratio* is a cross-condition R^2 divided by the target condition's within-condition R^2; the tautology, being common to both, is largely normalised out (it enters additively, so the cancellation is approximate, not exact). *Knee gap* is the raw same-λ SSM−baseline gap at λ_{L1}=1.

Table 4: Models, datasets, and source dictionaries. Confidence-filtered CounterFact sizes (70m: top-1,000 candidates by confidence — the optimiser-debug lane, excluded from the H1/H2 replication claim); SAE column gives the public suite used as frozen sources ("—" = diagnostics only). Seed-stability column is decoder match@0.7 across three fits; where a row names two dictionaries it lists them in the same order, and "—" means not measured (the Qwen2 trained-L1 fit is single-seed).

Model	L	D	N	SAE source (type)	seed stab.
Pythia-70m-deduped	6	512	1,000	EleutherAI 32k (TopK k=16)	0.99
Pythia-160m-deduped	12	768	767	MLSAE x64-k32 (TopK)	0.38
Pythia-410m-deduped	24	1024	1,555	MLSAE x64-k32 (TopK) / trained L1	0.45 / 1.00
GPT-2-small	12	768	1,358	res-jb (L1)	0.95
GPT-2-medium	24	1024	1,795	—	—
GPT-2-large	36	1280	1,924	—	—
Qwen2-0.5B	24	896	2,540	Chanin (2024) (L1, 12L prefix) / trained L1	0.81 / —

D Cross-model inventory

Table 4 lists the models, datasets, and source dictionaries.

Its seed-stability column is the cross-model half of the Section 6.5 finding: L1 sources stable on two additional families, TopK unstable on two Pythia models, with dataset size ruled out by subsampling GPT-2 to the smallest N (match 0.95 unchanged). The debug lane's EleutherAI TopK suite is nonetheless stable (0.99): TopK instability is not universal — every TopK source tested at $\geq$12 layers was unstable and every L1 source stable — which is why the controlled within-model contrast of Section 6.5, not the zoo, carries the mechanism claim.

E Reproducibility

All experiments ran on a single consumer GPU (NVIDIA GTX 1050, 3 GB, Pascal) with a pinned environment (Python 3.12; `torch` 2.5.1+cu124; `transformers` $<$ 5; TransformerLens 2.x); the constraint is documented because it shaped the design (activation caching, alive-set compression, fp16 loading for the larger models, and the rank-128 CLT being infeasible). Every run logs its resolved config, git commit, seed, and library versions before doing any work. Numeric intermediates (CSV/JSON tables) are tracked in git; heavy artefacts (activation stores, checkpoints) are regenerable from scripts. A single entry point — `scripts/reproduce.sh` — checks the environment, runs the test suite (208 tests), asserts dataset sizes, and regenerates all 16 paper figures byte-identically from the tracked tables; `-toy-pipeline` additionally re-runs the full Pythia-70m lane end to end. Code, protocol documents, and the progress tracker are released with the paper under a permissive licence.

Software: PyTorch (Paszke et al., 2019), TransformerLens (Nanda and Bloom, 2022), scikit-learn (Pedregosa et al., 2011), statsmodels (Seabold and Perktold, 2010), and the SAE suites cited in Section 5.2.